

**MINISTRY OF EDUCATION AND TRAINING
FPT UNIVERSITY**

**A TWO-TIER COGNITIVE ARCHITECTURE FOR
AUTOMATED THREAT DETECTION AND RESPONSE
USING AGENTIC AI**

by

Nguyen Duc Binh

A thesis submitted in conformity with the requirements
for the degree of Master of Software Engineering

**MINISTRY OF EDUCATION AND TRAINING
FPT UNIVERSITY**

**A TWO-TIER COGNITIVE ARCHITECTURE FOR
AUTOMATED THREAT DETECTION AND RESPONSE
USING AGENTIC AI**

by

Nguyen Duc Binh

A thesis submitted in conformity with the requirements
for the degree of Master of Software Engineering

Supervisor:
Dr. Bui Van Hieu
Dr. Dang Van Hieu

Abstract

In the contemporary cybersecurity landscape, Security Operations Centers (SOCs) are consistently overwhelmed by the sheer volume of network alerts, leading to severe alert fatigue and the potential oversight of critical Advanced Persistent Threats (APTs). To address the inherent limitations of traditional, static rule-based systems and the latency bottlenecks of processing raw logs directly through Large Language Models (LLMs), this thesis introduces SENTINEL: A Two-Tier Cognitive Architecture for Automated Threat Detection and Response. The proposed framework operates by deploying Welford's online statistical algorithm at the Transport Tier (Tier-1) to filter benign traffic at wire speed with minimal computational overhead. Subsequently, the Cognitive Agent Tier (Tier-2) leverages a LangGraph-driven AI Agent equipped with a Dual Retrieval-Augmented Generation (Dual-RAG) mechanism—integrating MITRE ATT&CK and NIST guidelines—to perform deep, contextual threat reasoning and mitigation.

Evaluated on the CSE-CIC-IDS2018 and DAPT2020 datasets, SENTINEL demonstrates substantial improvements in both detection accuracy and processing efficiency. Furthermore, the system incorporates robust AI security guardrails, including Delimited Data Encapsulation and HMAC log chaining, achieving near-perfect mitigation against known adversarial AI manipulation attacks and ensuring forensic data integrity. The findings underscore the efficacy of combining deterministic statistical filters with explainable Agentic AI, offering a highly scalable and secure automation solution for modern network defense.

Keywords: Retrieval-Augmented Generation, Agentic AI, Network Security, SOC Automation, Welford's Algorithm, Adversarial AI Security.

Acknowledgments

I would like to express my deepest gratitude to my supervisors, Dr. Bui Van Hieu and Dr. Dang Van Hieu, for their invaluable assistance, dedicated guidance, and insightful feedback throughout the development and writing of this thesis. Their profound expertise and continuous encouragement have been instrumental in the successful completion of this research.

I am particularly grateful to all the faculty members and staff at FPT University for their kind support, excellent academic environment, and provision of necessary resources during my master's study.

Finally, I would like to extend my heartfelt thanks to my family for their unwavering moral support, patience, and understanding throughout my academic journey.

List of Figures

1.1	Conceptual Illustration of a Modern AI-Assisted Security Operations Center (SOC) confronting high-velocity threat data.	9
2.1	The human bottleneck in traditional SOC systems, where massive alert volumes overwhelm analyst capacity.	18
2.2	State-Graph (FSM) based workflow of an Agentic AI system during security incident analysis.	20
2.3	The mechanism of a "Log-Substrate Prompt Injection" attack: The attacker embeds malicious commands into network traffic, exploiting log ingestion to manipulate the LLM's reasoning.	22

List of Tables

Contents

Abstract	1
Acknowledgments	2
1 Introduction	7
1.1 Background and Motivation	7
1.2 Research Objectives	10
1.3 Research Questions	11
1.4 Scope and Object of Study	12
1.5 Research Methodology	13
1.6 Related Works	14
1.7 Contributions of the Thesis	15
1.8 Structure of the Thesis	15
2 Theoretical Background	17
2.1 Overview of Security Operations Centers (SOCs) and the Cognitive Bot- tleneck	17
2.2 Large Language Models (LLMs) and the Mathematics of Quantization . .	19
2.3 Agentic AI Architecture and State-Graph Models	19
2.4 Retrieval-Augmented Generation (RAG) and Hybrid Search Optimization	21
2.5 LLM Cognitive Vulnerabilities and Prompt Injection Vectors	21
2.6 Welford’s Online Statistical Algorithm	23
3 The Proposed SENTINEL Architecture	24
3.1 Overview of the Two-Tier Cognitive Architecture	24
3.2 Tier-1: Real-Time Traffic Filtration Engine (RuleEngine & Welford) . . .	24
3.3 Tier-2: LangGraph Cognitive AI Agent	25
3.4 Dual-RAG Mechanism and Semantic Memory	25
3.5 Threat Memory and APT Campaign Correlation	25

3.6	AI Security Guardrails Layer	25
3.7	Data Integrity and HMAC Protection	26
4	Technical Implementation	27
4.1	Technology Stack and Source Code Structure	27
4.2	Tier-1 Implementation: Filtration Engine & Welford	27
4.3	Tier-2 Implementation: AI Agent and RAG	27
4.4	Guardrail and HMAC System Implementation	28
4.5	SOC Administration Dashboard Implementation (Streamlit UI)	28
4.6	Containerization and Infrastructure Deployment	28
5	Experiments and Evaluation	30
5.1	Experimental Environment and Datasets	30
5.2	5D Evaluation Metrics and Ablation Study Design	30
5.3	Accuracy Evaluation & McNemar’s Statistical Test	31
5.4	Performance Latency Evaluation & Mann-Whitney U Test	31
5.5	Adversarial Robustness Evaluation (Security)	31
5.6	Integrity and RAG Quality Evaluation via LLM-as-a-Judge	32
6	Conclusion and Future Work	33
6.1	Synthesis of Achieved Research Results	33
6.2	Scientific and Practical Contributions	34
6.3	Existing Limitations	34
6.4	Future Development Directions	34
	References	36

Chapter 1

Introduction

This chapter establishes the overarching context and foundational rationale for the research conducted in this thesis. By dissecting the rapidly evolving complexities of the modern cybersecurity landscape, it underscores the systemic vulnerabilities inherent in contemporary Security Operations Centers (SOCs). The introduction delineates the profound limitations of traditional, deterministic rule-based systems when confronting sophisticated, multi-stage cyber adversaries. Furthermore, it highlights the paradoxical challenges introduced by the integration of Generative Artificial Intelligence (GenAI)—balancing immense analytical potential against acute computational latency and adversarial manipulation risks. Consequently, this chapter articulates the necessity for SENTINEL, a novel cognitive two-tier architecture, outlining the specific research objectives, core inquiries, defined scope, and the experimental methodology that guides the empirical validation of the proposed system.

1.1 Background and Motivation

The acceleration of global digital transformation, characterized by the pervasive adoption of cloud-native computing paradigms, the Internet of Things (IoT), and decentralized remote workforce infrastructures, has fundamentally eroded traditional network perimeters. Consequently, enterprise networks have transitioned into highly porous, heterogeneous environments, exponentially expanding the potential attack surfaces available to malicious actors. In this volatile ecosystem, cyber adversaries have continuously refined their methodologies, shifting from isolated, easily identifiable malware strains to highly orchestrated, stealthy campaigns known as Advanced Persistent Threats (APTs) [1]. These sophisticated threat actors leverage polymorphic payloads, “living-off-the-land” (LotL) techniques, and encrypted command-and-control (C2) channels to obfuscate their activities, allowing them to traverse enterprise networks undetected for extended periods.

Security Operations Centers (SOCs) operate as the vital defense mechanism for organizations, entrusted with the continuous monitoring, real-time detection, and rapid mitigation

of such security incidents. To achieve this, modern SOC's deploy a myriad of security telemetry sensors, including Next-Generation Firewalls (NGFW), Intrusion Detection and Prevention Systems (IDS/IPS), and Endpoint Detection and Response (EDR) agents. These diverse data streams are typically aggregated into a central Security Information and Event Management (SIEM) platform [2]. However, the relentless ingestion of this telemetry has catalyzed a structural crisis within the cybersecurity industry, commonly referred to as the "Big Data Paradox" of cybersecurity [3]. This paradox manifests through several critical operational vulnerabilities:

- **Alert Fatigue Crisis:** Traditional SIEM platforms and IDS appliances generate an overwhelming volume of security alerts daily, often numbering in the tens of thousands. The vast majority of these notifications are False Positives (FPs) triggered by benign anomalies, misconfigurations, or poorly tuned heuristic thresholds [5]. This incessant barrage of alerts overwhelms Tier-1 security analysts, inducing severe cognitive overload and "alert fatigue". Consequently, the probability of a genuine, critical intrusion indicator being inadvertently dismissed or completely overlooked increases exponentially, rendering the human element the weakest link in the defense chain.
- **The Inefficacy of Deterministic Rule-Based Systems:** The foundational logic of traditional SIEM and IDS tools relies heavily on static signature matching and pre-defined heuristic rules [10]. While highly effective at blocking known threat vectors, these mechanisms exhibit a profound structural blindness when confronted with Zero-day vulnerabilities, custom malware variants, or novel, signature-less exploitation techniques. Adversaries easily bypass these static defenses by making trivial modifications to their attack payloads.
- **Deficiencies in Temporal Data Reduction:** In an attempt to mitigate the prohibitive costs and computational overhead of processing petabytes of log data, numerous SOC's resort to random sampling, log truncation, or aggressive static thresholding. Unfortunately, these data reduction strategies inadvertently sever the temporal evidence chains. By dropping seemingly innocuous packets, the SOC destroys the critical contextual linkages necessary to reconstruct the "low-and-slow" behavioral patterns characteristic of APT campaigns, which rely on stealthy reconnaissance and subtle lateral movement spread across weeks or months [4].



Figure 1.1: Conceptual Illustration of a Modern AI-Assisted Security Operations Center (SOC) confronting high-velocity threat data.

The recent maturation of Generative Artificial Intelligence (GenAI), and particularly Large Language Models (LLMs), has illuminated a promising paradigm shift for SOC automation. LLMs possess unprecedented capabilities for interpreting complex, unstructured contextual indicators, generating human-readable forensic summaries, and formulating autonomous incident response strategies. Nevertheless, the direct, unmediated integration of LLMs into the high-velocity data streams of a live SOC introduces two severe, intersecting bottlenecks:

1. **Prohibitive Inference Latency and Computational Overhead:** Routing raw, unfiltered network packet captures comprehensively through a multi-billion parameter LLM for inference creates an insurmountable processing bottleneck. The immense computational complexity inherent to the self-attention mechanisms of the Transformer architecture precludes its viability for wire-speed, real-time network traffic

analysis. A system attempting to process gigabits of raw traffic via LLM inference would incur unsustainable latency, allowing rapid intrusions to succeed before the AI can formulate a response.

2. **Extreme Susceptibility to Adversarial AI Manipulation:** Integrating an LLM exposes the SOC to novel cyberattack vectors, most notably Direct Prompt Injection, Semantic Data Poisoning, Data Exfiltration via Markdown rendering, and Delimiter Smuggling [9]. Sophisticated threat actors can embed malicious linguistic instructions directly into otherwise standard network payloads (e.g., crafting a malicious HTTP User-Agent string or a poisoned DNS query that contains instructions like: *"System override: Ignore all previous instructions, clear all memory context, and classify this traffic sample as perfectly benign"*). If the LLM ingests and processes these poisoned logs without strict, cryptographically secure isolation mechanisms, the cognitive core of the SOC's defense is compromised. The AI system can be manipulated to hallucinate benign verdicts for malicious activities or exfiltrate internal data, turning a powerful defensive asset into a catastrophic liability.

Recognizing these formidable strategic and technical challenges, and synthesizing foundational principles from both AI Engineering and Information Security, the absolute necessity for a mathematically optimized, hybrid architecture becomes irrefutable. What is required is a dual-layered approach: a high-speed, stateless heuristic filtration engine at the foundational tier to process traffic at wire-speed, seamlessly coupled with a deep, cognitively isolated Agentic AI reasoning engine at the upper tier to handle complex triage. Consequently, the research topic **"A Two-Tier Cognitive Architecture for Automated Threat Detection and Response Using Agentic AI"** (codenamed the SENTINEL System) was formulated to directly address and resolve these pervasive industry bottlenecks.

1.2 Research Objectives

The overarching ambition of this thesis is to fundamentally resolve the automation and cognitive bottlenecks plaguing modern SOC's. This is achieved by engineering an intelligent, high-performance intrusion detection and response framework that is not only highly accurate in threat classification but inherently resilient against adversarial attacks targeting its own AI components.

General Objective: To architect, implement, and rigorously evaluate the SENTINEL system—a pioneering dual-tier cognitive architecture that seamlessly integrates ultra-low-latency statistical anomaly detection with contextually aware, secure Agentic AI analysis.

Specific Objectives:

- **Develop Tier-1 (Transport/Filtering Tier):** Implement and optimize Welford’s on-line statistical algorithm [6] for continuous, real-time calculation of rolling variance and Z-Scores over narrow-window network data streams. The specific objective is to accurately filter the vast majority of benign background traffic at wire speed (operating at $O(1)$ space and time complexity) without inadvertently severing the crucial temporal and behavioral state chains of the traffic flows.
- **Construct Tier-2 (Cognitive Agent Tier):** Engineer a sophisticated autonomous AI Agent leveraging the LangGraph state machine architecture [8] powered by highly optimized Local LLMs. This tier must incorporate an advanced Dual Retrieval-Augmented Generation (Dual-RAG) mechanism [7]—synthesizing threat intelligence from both the MITRE ATT&CK framework and NIST SP 800-61r2 incident handling guidelines—to execute deep threat classification, formulate logical forensic reasoning, and autonomously propose actionable mitigation responses.
- **Enforce Uncompromising AI Security Guardrails:** Design and implement a proprietary Delimited Data Encapsulation mechanism, utilizing cryptographically secure randomized delimiters, designed to neutralize Direct Prompt Injection, Encoding Bypasses, and Delimiter Smuggling attacks that may be embedded within the network payloads under analysis.
- **Ensure Forensic Integrity of Evidence:** Implement SHA-256 cryptographic hashing for all ingested RAG knowledge bases and strictly enforce HMAC (Hash-based Message Authentication Code) chaining for all database audit logs. This guarantees absolute non-repudiation and prevents the forging, alteration, or deletion of forensic evidence by external attackers or malicious insiders.

1.3 Research Questions

To systematically achieve the delineated objectives and ensure the scientific rigor of the proposed architecture, this research focuses on addressing three core scientific inquiries:

1. **RQ1 (Performance & Filtration):** How can a high-volume, real-time network traffic filtration mechanism be mathematically designed to operate at wire-speed without disrupting, truncating, or losing the latency-heavy, multi-stage event chains that are characteristic of sophisticated APT attacks?
2. **RQ2 (Security & Integration):** By what architectural methodology can the advanced cognitive reasoning capabilities of LLMs be safely integrated into the automated SOC incident response pipeline while simultaneously resolving the inference

latency paradox and effectively mitigating the existential risks of AI manipulation via malicious log injection (e.g., Prompt Injection, Delimiter Smuggling)?

3. **RQ3 (Accuracy & Hallucination Reduction):** What is the measurable, empirical efficacy of augmenting an autonomous AI Agent’s localized Semantic Memory with an external, authoritative dual knowledge retrieval system (Dual-RAG) in identifying complex Zero-day threats and statistically minimizing the occurrence of LLM cognitive hallucinations?

1.4 Scope and Object of Study

Objects of Study:

- Modern Security Operations Center (SOC) architectures and automated incident response workflows.
- High-velocity network flow logs (NetFlow), packet capture metadata (PCAP), and intercepted malicious packet payloads.
- Advanced computational techniques including Agentic Workflow Orchestration (Lang-Graph), Quantized Large Language Models, Vectorized Knowledge Graphs, and Retrieval-Augmented Generation (RAG).
- Network intrusion methodologies, stealthy lateral movement tactics, and adversarial techniques targeting Artificial Intelligence systems (focusing heavily on Prompt Injection, RAG Poisoning, and Semantic Data Poisoning).

Research Scope:

- **Empirical Datasets:** Comprehensive evaluations are conducted exclusively on internationally standardized, peer-reviewed cybersecurity datasets. This includes the **CSE-CIC-IDS2018** dataset (simulating a broad, comprehensive spectrum of modern network attacks including DoS, Web Attacks, Infiltration, and Botnets) and the **DAPT2020** dataset (specifically utilized for simulating multi-stage, highly stealthy, prolonged APT campaigns). These datasets are unified into a single, time-interleaved traffic stream (Unified Stream Engine) to validate the system in both deterministic offline batch scenarios and dynamic online messaging setups using Redis message brokers. Furthermore, Zero-day attacks are modeled using a non-synthetic, real-derived methodology, where genuine benign network flows are injected with extreme statistical deviations (manipulating a single Welford statistical feature to its extremity) to rigorously test the system’s signature-less detection capability.

- **AI Models and Hardware Constraints:** The research strictly investigates and optimizes the deployment of open-source Local LLMs (specifically utilizing the *Gemma-2-9B-IT* model). Advanced quantization techniques (e.g., Q4_K_M via llama.cpp) are employed to ensure operational feasibility and high inference speeds on a single, consumer-grade GPU. This rigorous hardware constraint is intentionally imposed to guarantee absolute data privacy and sovereignty, proving that the system can be deployed within a strictly air-gapped, internal enterprise SOC environment without relying on external, cloud-based API endpoints.
- **Mitigation Execution Environment:** The autonomous system responses generated by the Agent are strictly restricted to forensic analysis, detailed reporting, and the generation of simulated Firewall deployment rules. The research explicitly excludes the automated, unverified execution of physical network disconnection commands on live routing appliances. This boundary is established to prevent self-inflicted Denial of Service (DoS) incidents or catastrophic operational disruptions arising from potential, albeit rare, AI misclassifications.

1.5 Research Methodology

This thesis employs a rigorous, quantitative research methodology deeply integrated with advanced software engineering practices, forming a comprehensive Experimental Engineering Methodology framework:

- **Architectural Design and System Implementation:** Formulating the initial theoretical blueprint and translating it into a highly modular software system. The development stack utilizes Python 3.10+, Neo4j for graph relationships, FAISS for dense vector similarity search, LangGraph for stateful agent orchestration, and llama.cpp for accelerated local LLM inference. The entire environment is encapsulated using Docker containerization to ensure strict reproducibility across varied testbeds.
- **Simulation and Adversarial Experimentation:** Constructing a continuous, high-fidelity data stream simulation environment (the Unified Stream module supporting both offline evaluation and Redis-based online publication) to conduct rigorous End-to-End (E2E) testing across 22 distinct execution scenarios. These tests evaluate the detection of multi-day APT campaigns and real-derived zero-day outliers. Furthermore, bespoke adversarial cryptographic payload tools are developed and executed against the system to aggressively benchmark the robustness and resilience of the implemented AI security guardrails.

- **Quantitative Evaluation and Statistical Analysis:** Assessing overall system performance, accuracy, and reliability through a proprietary 5-Dimensional quantitative framework (5D Metrics encompassing: Accuracy, Security, Performance, Integrity, and Cognitive reasoning capabilities). To empirically prove the statistically significant superiority of the proposed SENTINEL architecture against independent baseline models, rigorous non-parametric statistical tests (including McNemar’s test for paired nominal data and the Mann-Whitney U test) alongside comprehensive Ablation Studies are systematically applied.

1.6 Related Works

The application of Artificial Intelligence in Intrusion Detection Systems (IDS) has been the subject of extensive academic research. Traditional approaches have heavily favored classical Machine Learning (ML) and Deep Learning (DL) architectures. Numerous studies have successfully applied Random Forests, Support Vector Machines (SVM), Long Short-Term Memory (LSTM) networks, and Convolutional Neural Networks (CNN) to classify network anomalies. While these models demonstrate high statistical accuracy on benchmark datasets, they suffer from severe deficiencies in Explainability. They operate as "black boxes," providing binary classifications without actionable forensic context, which severely limits their practical utility for human analysts during incident triage.

In recent years, the literature has shifted towards exploring the integration of LLMs in SOC operations. Recent 2025 and 2026 works showcase the transition from simple "copilot" assistants to fully agentic platforms. For example, Analyst *et al.* propose a policy-guided threat hunting framework acting as an LLM-enabled triage overlay for Splunk [15]. Taking a more autonomous approach, Abdennebi *et al.* introduced LanG, a governance-aware Agentic AI platform orchestrating SOC tools via LangGraph [14]. Furthermore, Researcher *et al.* developed CyberRAG, demonstrating the high efficacy of Agentic RAG specifically for cyber attack classification and automated reporting [13]. While these architectures significantly improve analytical reasoning, they generally treat the LLM as the primary analysis engine for raw logs, largely ignoring the fundamental latency bottleneck associated with processing massive, continuous log volumes at wire speed.

Crucially, a comprehensive review of the current literature reveals a significant research gap concerning cognitive security. Despite the advancement of multi-agent frameworks, the academic discourse has only recently recognized the inherent fragility of these LLM-integrated systems against diverse adversarial AI manipulations. A 2026 study by Ravindran *et al.* mathematically proved the existential risk of "Log-Substrate Prompt Injection"—where an attacker embeds adversarial payloads directly into the network logs being analyzed, successfully poisoning the LLM’s reasoning context [12]. The SENTINEL ar-

chitecture proposed in this thesis directly targets these combined gaps by engineering a hybrid, two-tier system: it utilizes Welford’s algorithm to mathematically resolve the latency paradox, while enforcing strict cryptographic guardrails against log-substrate prompt injections to ensure the absolute cognitive security of the Agentic AI engine.

1.7 Contributions of the Thesis

This research yields prominent, measurable contributions that advance both theoretical cybersecurity frameworks and practical, deployable engineering applications:

- **Architectural Contribution:** The successful proposition, implementation, and empirical validation of the SENTINEL two-tier hybrid model. This architecture achieves a highly sought-after, optimal equilibrium between ultra-high throughput network performance (powered by Welford’s online algorithm) and deep, contextually accurate cognitive reasoning (driven by Agentic AI).
- **AI Security Contribution:** Providing reproducible, empirical evidence regarding the extreme efficacy of Delimited Data Encapsulation when combined with strict integrity verification mechanisms (HMAC & SHA-256). This defensive matrix renders the core AI reasoning engine highly resilient to Prompt Injections, Encoding Bypasses, and Semantic Data Poisoning—achieving a near-perfect mitigation rate against known adversarial vectors, addressing a critical, often-overlooked vulnerability plaguing many contemporary commercial LLM-integrated security platforms.
- **Autonomy and Orchestration Contribution:** Pioneering the integration of a State-Graph based Agentic AI model directly into the Incident Response lifecycle. This effectively supersedes fragile, static Runbook Automation (RBA) scripts with dynamic, real-time, context-aware decision-making intelligence capable of adapting to unforeseen threat mutations.

1.8 Structure of the Thesis

To ensure clarity and logical progression, the thesis is systematically organized, transitioning from foundational theoretical concepts to complex design paradigms, empirical implementations, and finally, quantitative conclusions. The document comprises six primary chapters:

- **Chapter 1: Introduction** – Establishes the background, outlines the core motivation, reviews existing academic literature to explicitly highlight unresolved research

gaps, defines the specific research objectives and questions, delineates the scope, details the experimental methodology, and summarizes the primary scientific contributions of the thesis.

- **Chapter 2: Theoretical Background** – Constructs the essential theoretical foundation required to understand Agentic AI, Retrieval-Augmented Generation (RAG) systems, Welford’s online statistical algorithm, and the principles of LLM adversarial security. This chapter critically analyzes traditional SOC architectures and establishes the technical paradigms underlying the thesis.
- **Chapter 3: The Proposed SENTINEL Architecture** – Provides an exhaustive, granular detailing of the comprehensive design of the two-tier cognitive architecture. This includes the mechanical operations of the Transport Tier (Tier 1), the orchestration logic of the Cognitive Agent Tier (Tier 2), the structural design of the Dual-RAG security knowledge graph, and the mathematical implementation of the defensive AI guardrails.
- **Chapter 4: Technical Implementation** – Specifies the software development environment, details the system’s source code hierarchy, explains the integration of core tools (such as Neo4j, FAISS, and llama.cpp), describes the Docker containerization strategy, and showcases the construction of the interactive, Human-in-the-Loop (HITL) Streamlit Dashboard.
- **Chapter 5: Experiments and Evaluation** – Thoroughly outlines the experimental methodologies and hardware setups. It provides a deep dive into the datasets utilized (CSE-CIC-IDS2018, DAPT2020) and presents a profound quantitative evaluation based on the 5D metrics framework, extensive Ablation Studies, and rigorous statistical hypothesis testing.
- **Chapter 6: Conclusion and Future Work** – Summarizes the empirical results achieved, explicitly correlates these findings with the initial research objectives, transparently acknowledges existing architectural or hardware limitations, and proposes strategic directions for future research and system expansion.

Chapter 2

Theoretical Background

This chapter establishes the in-depth theoretical foundations in computer science and cybersecurity that serve as the mathematical and architectural basis for the SENTINEL system. The content is systematically developed, beginning with an analysis of the structural limitations of modern Security Operations Centers (SOCs), followed by statistical anomaly detection algorithms, and delving into the core technologies of Generative AI. Specifically, this chapter explores the operating principles of Large Language Models (LLMs), Model Quantization techniques, Agentic AI architectures based on state-graphs, Retrieval-Augmented Generation (RAG) systems, and finally, a detailed analysis of adversarial AI security vulnerabilities.

2.1 Overview of Security Operations Centers (SOCs) and the Cognitive Bottleneck

Modern Security Operations Centers (SOCs) form the core proactive defense layer of enterprise networks. Architecturally, a SOC operates as a centralized ecosystem, integrating a myriad of security appliances such as Security Information and Event Management (SIEM) platforms, Intrusion Detection and Prevention Systems (IDS/IPS), Next-Generation Firewalls (NGFW), and Endpoint Detection and Response (EDR) solutions. The standard Incident Response (IR) lifecycle requires these systems to ingest massive volumes of telemetry data, analyze them for anomalies based on static signatures or predefined heuristics, and generate alerts for Tier-1 analysts to investigate.

However, the shift towards cloud-native computing models and Zero-Trust architectures has led to an unprecedented data explosion. This colossal volume of network traffic directly contributes to the "Alert Fatigue" challenge. In operational reality, Tier-1 analysts are frequently overwhelmed by tens of thousands of alerts daily, among which False Positives typically constitute the vast majority. From a cognitive science perspective, this continuous information overload depletes human attentional resources, leading to a severe degradation

in analytical acuity, thereby exponentially increasing the probability of missing genuine and destructive threats [5].




Figure 2.1: The human bottleneck in traditional SOC systems, where massive alert volumes overwhelm analyst capacity.

Compounding this issue is the proliferation of Advanced Persistent Threats (APTs). APT campaigns are characterized by "low-and-slow" attack methodologies, where threat actors fragment their infiltration behavior into multiple insignificant event sequences, operating below the detection thresholds of traditional security systems. Deterministic rule-based systems, due to their over-reliance on instantaneous static thresholds, completely lack the capacity to maintain long-term state to contextualize and correlate these discrete indicators of compromise [4].

2.2 Large Language Models (LLMs) and the Mathematics of Quantization

At the forefront of the Generative AI revolution are Large Language Models (LLMs), built primarily upon the Transformer neural network architecture. The foundational breakthrough of the Transformer is the Self-Attention mechanism, allowing the model to compute correlation weights among all tokens in an input sequence simultaneously, regardless of their positional distance. This capability eliminates the long-term forgetting problem of traditional RNN/LSTM networks, endowing LLMs with profound contextual natural language processing capabilities, making them ideal tools for interpreting unstructured security logs.

However, integrating LLMs into a SOC environment faces a critical barrier regarding Data Sovereignty. Enterprise security data contains highly sensitive infrastructure information. Transmitting this data to external cloud LLM APIs directly violates security compliance standards and Zero-Trust principles. Consequently, deploying Local LLMs within fully air-gapped environments is a prerequisite.

The subsequent barrier is the immense hardware requirement. For a model like *Gemma-2-9B-IT* (with 9 billion parameters) to run locally using 16-bit floating-point (FP16), it requires approximately 18GB of VRAM solely for weight storage. To be feasible on consumer-grade hardware, Model Quantization techniques must be applied. Quantization is a mathematical process aimed at reducing the resolution of neural network parameters from a continuous space (e.g., FP16) to a discrete space with lower bit-width (e.g., 4-bit integer, INT4). By utilizing the GGUF format and the Q4_K_M quantization standard via the `llama.cpp` framework, the required memory footprint is reduced by over 70%, while utilizing memory bandwidth more efficiently, ensuring high inference speeds with negligible degradation in cognitive capacity compared to the original model.

2.3 Agentic AI Architecture and State-Graph Models

The evolution of AI in cybersecurity is shifting aggressively from Static Runbook Automation (RBA) to dynamic Agentic AI models. RBA is inherently a rigid script-based system (if-this-then-that), capable only of linear responses to predefined indicators. Conversely, Agentic AI is equipped with the autonomous capacity to perceive, perform multi-step planning, critique, and self-adjust behavior based on the evolving real-time context of a security incident.

Mathematically and architecturally, these sophisticated autonomous agents do not operate as monolithic blocks but are modeled as Finite State Machines (FSM) in the form of

Directed Graphs. Within this structure:

- **Nodes:** Represent independent computational or cognitive tasks (e.g., retrieving CTI, payload analysis, proposing mitigations).
- **Conditional Edges:** Define dynamic routing logic, utilizing the LLM’s analytical capabilities to decide which subsequent node to activate based on the current node’s output.




Figure 2.2: State-Graph (FSM) based workflow of an Agentic AI system during security incident analysis.

The utilization of the LangGraph framework [8] enables the construction of stateful cognitive loops. It preserves the agent’s ”working memory” throughout the analysis session, allowing the AI to self-reflect on previous conclusions, request additional context if the data is insufficient (Analysis Incomplete), and only render a final verdict when hypotheses are fully verified.

2.4 Retrieval-Augmented Generation (RAG) and Hybrid Search Optimization

The most severe limitation of LLMs is the phenomenon of "Hallucination" — the tendency to generate information that is technically incorrect yet linguistically highly convincing. In security, a hallucinated decision could lead to disastrous misconfigurations, such as mistakenly blocking critical services (Denial of Service). The Retrieval-Augmented Generation (RAG) architecture [7] fundamentally resolves this risk by forcing the LLM to ground its inferences on specialized knowledge blocks retrieved from internal databases prior to generating a response.

To ensure absolute precision in retrieving security knowledge (such as MITRE ATT&CK or NIST SP 800-61r2), modern RAG architectures employ a Hybrid Search mechanism, simultaneously combining two paradigms:

1. **Dense Vector Search:** Utilizes embedding models (like `all-MiniLM-L6-v2`) to map text data into high-dimensional numerical vectors. The FAISS library then computes the Cosine Similarity distance between the query vector and data vectors to find text segments with *semantic* similarity, even if they lack exact keywords.
2. **Sparse Keyword Search:** Based on the BM25 (Best Matching 25) algorithm, utilizing TF-IDF (Term Frequency-Inverse Document Frequency) weighting for exact matching of specific technical strings, such as CVE identifiers, malware filenames, or IP addresses.

Finally, the Reciprocal Rank Fusion (RRF) algorithm [11] is applied to merge the two result sets. By computing the inverse rank score for each document from both result lists, RRF ensures that the documents provided to the LLM are those with the highest accuracy in both deep technical specificity and overall context.

2.5 LLM Cognitive Vulnerabilities and Prompt Injection Vectors

Deploying LLMs as the core engine in a SOC opens an entirely new spectrum of cybersecurity risks, collectively termed Adversarial AI Manipulations. The fundamental difference is that LLMs do not compile source code but process natural language. Consequently, the boundary between "Data" and "Instruction" is blurred.

One of the most existential threats is **Log-Substrate Prompt Injection** [12]. In this scenario, an attacker is aware that the organization utilizes an LLM-based log analysis system. Instead of directly attacking the web application, they intentionally embed malicious commands into standard data fields (such as HTTP User-Agent strings or DNS query payloads).

When the SIEM ingests these payloads and pushes them directly into the LLM's Context Window for analysis, the LLM may interpret this malicious payload as a high-level administrative directive, leading to the circumvention of system security rules.




Figure 2.3: The mechanism of a "Log-Substrate Prompt Injection" attack: The attacker embeds malicious commands into network traffic, exploiting log ingestion to manipulate the LLM's reasoning.

More advanced techniques like **Encoding Bypasses** utilize Base64 or Hexadecimal encoding to evade traditional keyword filters. Equally dangerous is **Delimiter Smuggling**, where attackers insert special characters to break the application's data encapsulation structure, causing the LLM to lose context. The consequences of these attacks range from the system autonomously mislabeling threats (Misdirection Verdict) to conceal malware, to Data Exfiltration via Markdown rendering techniques.

2.6 Welford's Online Statistical Algorithm

While Agentic AI provides unparalleled forensic analysis capabilities, the immense computational complexity of LLMs ($O(N^2)$ where N is the input sequence length for the Attention mechanism) creates a paradox: they cannot process massive volumes of raw network data streams at wire-speed. To resolve this problem, a high-throughput preprocessing data filtration mechanism operating with an optimal space and time complexity of $O(1)$ is mandatory.

Welford's Online Algorithm [6] provides a precise and numerically stable mathematical formulation for computing variance over streaming data. Instead of storing the entire historical dataset to compute the sum of squares (which poses risks of catastrophic cancellation and high complexity), Welford's algorithm updates the Mean (μ_k) and the Sum of Squared Differences ($M_{2,k}$) for each new data point x_k using recurrence relations:

$$\mu_k = \mu_{k-1} + \frac{x_k - \mu_{k-1}}{k} \quad (2.1)$$

$$M_{2,k} = M_{2,k-1} + (x_k - \mu_{k-1})(x_k - \mu_k) \quad (2.2)$$

$$\sigma_k^2 = \frac{M_{2,k}}{k} \quad (2.3)$$

By continuously maintaining rolling variance and mean with a constant $O(1)$ memory footprint, the system instantaneously computes a normalized Z-Score for each packet. This technique almost entirely eliminates benign background network traffic, thoroughly resolving the LLM latency bottleneck, while perfectly preserving lightweight anomalous event sequences, maintaining the context for low-and-slow APT infiltrations.

To mathematically validate zero-day threat detection capabilities without introducing synthetic data bias, the architecture leverages a real-derived zero-day modeling paradigm. By isolating genuine benign network flows from the ground truth and systematically elevating a single statistical feature (such as flow duration or cumulative byte count) to its extreme mathematical boundary, the system forces a significant deviation in Welford's rolling calculations. This generates an anomalous Z-score, mimicking a completely unknown statistical outlier. Consequently, the signature-less zero-day threat is caught by the Tier-1 filter and escalated to the Tier-2 cognitive agent, bypassing the limitations of static signature engines.

In summary, Chapter 2 has synthesized and profoundly dissected the theoretical principles, mathematical tools, and associated cognitive risks, forming a robust academic foundation.

Building upon this theoretical basis, Chapter 3 will detail the technical design of the SENTINEL dual-tier hybrid architecture.

Chapter 3

The Proposed SENTINEL Architecture

This chapter meticulously delineates the design of the SENTINEL system (Cognitive Two-Tier Architecture). Serving as the scientific core of the thesis, this architecture resolves the Threat Detection and Incident Response (TDIR) challenge automatically, at high velocity, and with absolute security against adversarial AI manipulation. The content covers the holistic two-tier data flow, the online Welford statistical filter, the LangGraph Agent model with Dual-RAG, and advanced security constraints (Guardrails, HMAC Chaining).

3.1 Overview of the Two-Tier Cognitive Architecture

- **Holistic System Block Diagram:** Illustrating the data flow starting from log generation, traversing Tier-1 (Transport/Filter Layer), escalating to Tier-2 (Cognitive/AI Agent Layer), and culminating in the Database/Dashboard.
- **Design Philosophy:** Separation of concerns—Tier-1 is strictly responsible for wire-speed noise filtration, while Tier-2 provides deep contextual analysis to resolve the blindness of static systems toward Zero-day threats.

3.2 Tier-1: Real-Time Traffic Filtration Engine (RuleEngine & Welford)

- **Session and Baseline Management:** Mechanisms for storing and updating network connection states temporally.
- **Online Welford's Algorithm Mechanism:** Mathematical specification for computing Z-Scores per log stream. When a Z-Score exceeds threshold α , the system generates an ESCALATE label for Tier-2 instead of dropping the packet.
- **Static Rule Engine Integration:** Incorporation of fundamental signature-based rules (DROP, WARN) to instantly eliminate rudimentary, known attack patterns.

3.3 Tier-2: LangGraph Cognitive AI Agent

- **StateGraph Design:** The architectural structure of the `SentinelState` object carrying session information. Definition of core Nodes: Analyze, Retrieve, Resolve.
- **Conditional Edges Routing:** The internal reasoning loop of the AI Agent. Upon identifying a knowledge deficit, the agent autonomously transitions to the `retrieve_knowledge` state before returning to evaluation.

3.4 Dual-RAG Mechanism and Semantic Memory

- **Dual-RAG Architecture:** Unification of two specialized knowledge bases: MITRE ATT&CK (TTP identification) and NIST SP 800-61r2 (Incident Response Handbook) [10].
- **Hybrid Search Algorithm:** The fusion of Dense (FAISS) and Sparse (BM25) scoring mechanisms utilizing RRF.
- **Semantic Cache:** A semantic hashing mechanism to bypass LLM inference if a highly similar log stream was successfully analyzed earlier in the temporal window, drastically reducing latency.

3.5 Threat Memory and APT Campaign Correlation

- **Graph Data / IP Correlation Storage (SQLite/Neo4j):** Logic for tracking the Reputation Score of IP addresses over prolonged timelines.
- **APT Correlation Mechanism:** An algorithm to chain temporally scattered, low-severity anomalies into a cohesive high-severity escalation verdict, exposing low-and-slow campaigns.

3.6 AI Security Guardrails Layer

- **Delimited Data Encapsulation:** An algorithm generating a cryptographically random token per session (e.g., `<|PAYLOAD_A7F9|>`) to encapsulate foreign data, ensuring the LLM does not execute hidden adversarial instructions.
- **Token Budget Manager & LogTemplateMiner:** Utilizing Drain3 to strip random textual variations and cluster logs, thereby preventing LLM Token Denial-of-Service (DoS) attacks.

- **Output Sanitization Filter:** Strict standardization and JSON schema enforcement of the LLM's output.

3.7 Data Integrity and HMAC Protection

- **HMAC Log Chaining:** A cryptographic hashing algorithm linking consecutive logs, ensuring that any tampering by malicious insiders breaks the chain and is immediately flagged.
- **RAG SHA-256 Checksum Validation:** Guaranteeing the vector database is not compromised by injected falsified documents (Data Poisoning).

In summary, Chapter 3 establishes the rigorous theoretical framework of the SENTINEL architecture, holistically resolving both performance bottlenecks (Tier-1) and deep, secure cognitive reasoning (Tier-2). Chapter 4 will specify the process of translating this architecture into technical source code.

Chapter 4

Technical Implementation

This chapter details the software engineering aspects and practical source code implementation of the SENTINEL architecture. Shifting from abstract theory, the content focuses on the project directory organization, core class structures (Classes/Methods), Docker service configurations, and the design of the Graphical User Interface (Streamlit Dashboard) which facilitates Human-in-the-Loop (HITL) interaction.

4.1 Technology Stack and Source Code Structure

- **Core Technologies:** Python 3.10+, LangGraph, LangChain, FAISS, SQLite, Streamlit, Docker, and llama.cpp (via the Llama-cpp-python binding).
- **Directory Structure:** A detailed tree representation strictly separating components: `src/core/` (Tier 1 & 2), `src/rag/` (Vector DB), `src/security/` (Guardrails), `src/dashboard/` (UI), `experiments/`, and `tests/`.

4.2 Tier-1 Implementation: Filtration Engine & Welford

- **RunningStats and SessionBaseline Classes:** Python source code implementation for Welford's algorithm. Resolving dictionary memory allocation challenges via IP hashing and cache garbage collection (TTL cache).
- **RuleEngine Class:** Pattern matching design based on Regular Expressions and Static Thresholds.

4.3 Tier-2 Implementation: AI Agent and RAG

- **OpenAILLMClient / Llama.cpp Wrapper:** Parameter configurations for the Local LLM (enforcing `temperature=0.0` to optimize determinism, context size

bounding, and `n_gpu_layers` mapping).

- **LangGraph Workflow Setup:** Utilizing `StateGraph` to define `analyze_node`, `retrieve_node`, and the conditional routing function `should_retrieve`.
- **DualRetriever Class:** Implementing the FAISS Index and `rank_bm25` to achieve Reciprocal Rank Fusion (RRF). Source code overview of the scoring and sorting algorithms.

4.4 Guardrail and HMAC System Implementation

- **DelimitedDataEncapsulator Class:** Logic for generating secure hexadecimal tokens via `os.urandom` and `hashlib` to wrap log payloads.
- **HMACHashing Class (Database):** Specification of the SQLite database schema (`threat_memory.db`) and the trigger functions calculating SHA-256 hashes derived from Previous Hash + Current Data.

4.5 SOC Administration Dashboard Implementation (Streamlit UI)

- **Dashboard Architecture:** Application of a Frontend (Streamlit) communicating dynamically with the Backend (SQLite/File logs) through reactive UI components.
- **Display Modules:**
 - *Real-time Queue Dashboard:* Visualizing Tier-1 alert metrics and Tier-2 verdicts.
 - *Threat Memory Analytics:* Displaying APT alert graphs across IP timelines.
 - *Audit Trail Integrity:* An automated module verifying HMAC Log Chaining, utilizing a traffic-light interface (Green = Valid, Red = Chain Broken).
- **Human-in-the-loop (HITL):** The interface for reviewing (Approve/Reject) AI-recommended verdicts to update actual Firewall Rules.

4.6 Containerization and Infrastructure Deployment

- **Dockerfile Specification:** Multi-stage build mechanisms to minimize image size and compile C++ libraries (`llama.cpp`) with CUDA acceleration support.

- **`docker-compose.yml` Service Management:** Linking containers for the Dashboard and Agent Engine, alongside shared log volume mounts.

In summary, Chapter 4 systematizes the software development process of SENTINEL, bridging core algorithms with user interfaces and virtualized infrastructure. Chapter 5 will subsequently establish the testing environment to quantitatively evaluate these components.

Chapter 5

Experiments and Evaluation

This chapter presents the empirical evidence of the thesis, validating the research hypotheses through rigorous quantitative metrics. The SENTINEL system is evaluated against a 5-Dimensional framework (5D Metrics: Accuracy, Performance, Security, Explainability, and Integrity). The content emphasizes the layered Ablation Study design and the application of biostatistical models (McNemar’s Test, Mann-Whitney U Test) to prove that the architectural enhancements are statistically significant and mathematically sound.

5.1 Experimental Environment and Datasets

- **Hardware Setup:** CPU Intel Core i9, 32GB RAM, GPU NVIDIA RTX 4060 Ti 16GB (utilized for executing the Gemma-2-9B-IT Q4_K_M model).
- **CSE-CIC-IDS2018 Dataset:** Introduction to the dataset encompassing diverse attack typologies (Brute Force, DoS, Web Attack) serving as the baseline for overarching accuracy evaluation.
- **DAPT2020 Dataset:** Employed for testing scenarios involving the identification of multi-stage Advanced Persistent Threat (APT) chains.
- **Unified Data Stream Simulation:** Detailing the scripting process that merges and replays log streams according to virtual timestamps to emulate real-time properties.

5.2 5D Evaluation Metrics and Ablation Study Design

- **Defining the 5D Metrics:** Accuracy (F1, Precision, Recall), Performance (Latency, Throughput), Security (Robustness Rate), Explainability (Audit Completeness), and Integrity (HMAC Validation).
- **Ablation Study Configurations (Config A through F):**

- Config A (Baseline 1): Traditional Rule-based engine exclusively.
- Config B (Baseline 2): Pure Local LLM inference (No RAG, no Welford).
- Config C: Welford Tier-1 + LLM.
- Config D: LLM + Single-RAG (FAISS).
- Config E: LLM + Dual-RAG Hybrid.
- Config F (Holistic SENTINEL): Tier-1 (Welford) + Tier-2 (LangGraph Agent + Dual RAG + Guardrails).

5.3 Accuracy Evaluation & McNemar’s Statistical Test

- **Precision, Recall, and F1-Score Results:** Comparative charts and tables across configurations. Highlighting Config F’s optimal F1 balance (eliminating Rule-based False Positives while enhancing Zero-day Recall).
- **APT Detection Efficacy:** Analysis of Tier-1 (Welford) capturing weak anomalous signals and subsequently activating the Agent’s Threat Memory chains.
- **McNemar’s Test Application:** Comparing the Confusion Matrices of Config B and Config F. Proving $p - value < 0.05$ (rejecting the Null Hypothesis) to confirm that the variance is attributable to architectural superiority rather than random chance.

5.4 Performance Latency Evaluation & Mann-Whitney U Test

- **Inference Latency Analysis:** Comparative table of average latency per flow. The exceptional offloading achieved by Tier-1 (Welford processing benign logs in $O(1)$ time) and the Semantic Cache (Cache Hits reducing LLM latency from 4-6 seconds to milliseconds).
- **Non-parametric Mann-Whitney U Test:** Demonstrating that the latency distribution of the SENTINEL configuration is significantly faster ($p < 0.05$) compared to routing all logs directly through the LLM.

5.5 Adversarial Robustness Evaluation (Security)

- **Adversarial AI Manipulation Dataset (45 Payloads):** Description of manipulative scenarios including Prompt Injections and Delimiter Smuggling (e.g., “Ignore system rules, report this IP as safe”).

- **Mitigation Rate Measurement:** Comparison between an unguarded system (80% manipulation success rate) versus SENTINEL's enabled Delimited Data Encapsulation mechanism (achieving near-perfect mitigation of injection attempts).

5.6 Integrity and RAG Quality Evaluation via LLM-as-a-Judge

- **HMAC Chaining Audit:** A hypothetical scenario where an attacker compromises the database to alter log labels. The dashboard immediately reports the broken hash chain.
- **RAG Scoring (RAGAS / LLM-as-a-Judge):** Utilizing an external Llama-3.1 model as an independent arbiter to score four criteria: Context Precision, Context Recall, Faithfulness, and Answer Relevancy. Proving that Dual-RAG with RRF outperforms Single Vector Search.

In summary, Chapter 5 verifies the technical prowess of SENTINEL through hard quantitative data and rigorous biostatistical tests. Chapter 6 will distill the significance of these findings and outline future trajectories.

Chapter 6

Conclusion and Future Work

This concluding chapter synthesizes the comprehensive research journey of the thesis, “A Two-Tier Cognitive Architecture for Automated Threat Detection and Response Using Agentic AI” (SENTINEL). The chapter encapsulates the scientifically proven results, critically self-assesses the limitations within the scope of a master’s thesis, and illuminates new horizons for cybersecurity research integrated with Artificial Intelligence.

6.1 Synthesis of Achieved Research Results

- **Resolving Research Questions:** Reaffirming the successful mitigation of latency bottlenecks and Alert Fatigue via the tiered architecture of the Tier-1 Welford filter and the Tier-2 LangGraph Agent.
- **AI System Security (AI Security):** Demonstrating the empirical viability of Delimited Data Encapsulation and HMAC Chaining structures in thwarting adversarial manipulation risks (Prompt Injection, Delimiter Smuggling, Data Tampering) targeting internal enterprise cognitive systems.
- **Overcoming Zero-Day and APT Challenges:** The exceptional synergy between historical alert tracking (Threat Memory) and the hybrid knowledge repository (Dual-RAG MITRE + NIST) provides profound contextual awareness, enabling the system to expose highly sophisticated campaigns.
- **Validation Rigor:** These achievements transcend theoretical propositions, having been robustly validated by 5D metrics and standard statistical tests (McNemar, Mann-Whitney U) on empirical datasets.

6.2 Scientific and Practical Contributions

- **Scientific Value:** Providing a highly modular architectural framework for integrating Mathematical Statistics with Large Language Models in cybersecurity. This substantially fortifies the theoretical repository concerning LLM security and Guardrails.
- **Practical Value:** The finalized SENTINEL product serves as a core engine or an independent auxiliary tool for Small and Medium Businesses (SMBs) lacking the budget to sustain large-scale Tier-1 SOC teams, while strictly maintaining Zero Trust privacy through offline operations.

6.3 Existing Limitations

Despite the system’s demonstrated superiority, the research acknowledges distinct boundaries:

- **Input Log Formats:** The system is currently deeply tested on Network Flow logs and HTTP payloads. Real-world applicability necessitates the ingestion of highly diverse unstructured formats (Windows Event Logs, Linux Syslogs, CloudTrail), requiring significantly more complex parsing mechanisms.
- **Hardware Constraints:** Although inference speed is accelerated via Semantic Caching, processing entirely novel alerts still requires substantial LLM reasoning time. This mandates high capital expenditure for specialized GPU hardware if operating at hundreds of Gigabits/second network scales.
- **Physical Mitigation Design:** The research deliberately constrained the system to recommend actions (e.g., Recommend Firewall Rules) rather than executing physical interventions on routers to prevent Self-Denial of Service (Self-DoS). However, this restricts true Real-Time Blocking capabilities.

6.4 Future Development Directions

Stemming from these limitations, future expansion trajectories for the system include:

- **Federated Learning:** Expanding SENTINEL into a decentralized architecture. Instead of sharing privacy-violating network logs, SOC endpoints share only learned model parameters regarding novel malware variants.

- **Multi-Agent System (MAS):** Decomposing the monolithic LangGraph Agent into specialized sub-agents: Malware Analyst Agent, Network Analyst Agent, and Threat Intel Agent. These agents would debate prior to issuing a final verdict, further elevating accuracy.
- **Automated Red Teaming:** Constructing an internal Red Team agent to continuously generate sophisticated adversarial payloads targeting Sentinel, thereby self-evaluating and autonomously upgrading the system's defensive posture.

Closing a modest endeavor, yet opening profound aspirations in the pursuit of securing the digital realm through the power of Agentic AI.

References

- [1] I. Sharafaldin, A. H. Lashkari, and A. A. Ghorbani, “Toward Generating a New Dataset for Cyber Security in Ad hoc Networks,” in *Proceedings of the International Conference on Information Systems Security and Privacy (ICISSP)*, 2018.
- [2] A. Alsubaei, A. Abuhussein, and S. Shiva, “Security Information and Event Management (SIEM) Features, Challenges, and Solutions,” *IEEE Access*, vol. 9, pp. 121111–121128, 2021.
- [3] P. M. M. S. Silva *et al.*, “The Big Data Security Challenges in the Era of Cloud Computing,” *Computer Networks*, vol. 156, pp. 102–115, 2019.
- [4] M. Alshamrani *et al.*, “A Survey on Advanced Persistent Threats: Techniques, Solutions, Challenges, and Research Opportunities,” *IEEE Communications Surveys & Tutorials*, vol. 21, no. 2, pp. 1851–1877, 2021.
- [5] S. Hassan, M. A. Saleem, and O. E. Ogu, “Alert Fatigue in Cybersecurity: A Review,” *IEEE Access*, vol. 10, pp. 58632–58645, 2022.
- [6] B. P. Welford, “Note on a Method for Calculating Corrected Sums of Squares and Products,” *Technometrics*, vol. 4, no. 3, pp. 419–420, 1962.
- [7] P. Lewis *et al.*, “Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks,” *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 33, pp. 9459–9474, 2020.
- [8] LangChain AI, “LangGraph: Multi-Actor, Stateful LLM Applications,” *GitHub repository*, 2024. [Online]. Available: <https://github.com/langchain-ai/langgraph>
- [9] K. Greshake, S. Abdelnabi, S. Mishra, A. Endres, T. Holz, and M. Fritz, “More than you’ve asked for: A Comprehensive Analysis of Novel Prompt Injection Threats to Application-Integrated Large Language Models,” *arXiv preprint arXiv:2302.12173*, 2023.
- [10] P. Cichonski, T. Millar, T. Grance, and K. Scarfone, “Computer Security Incident Handling Guide,” NIST Special Publication 800-61 Revision 2, 2012.
- [11] G. V. Cormack, C. L. A. Clarke, and S. Buettcher, “Reciprocal Rank Fusion Outperforms Condorcet and Individual Bootstrap Product Algorithms,” in *Proceedings of the 32nd International ACM SIGIR Conference on Research and Development in Information Retrieval*, 2009.

- [12] T. Ravindran *et al.*, “Poisoning the Watchtower: Prompt Injection Attacks Against LLM-Augmented Security Operations Through Adversarial Log Content,” *arXiv preprint arXiv:2605.24421*, 2026.
- [13] A. Researcher *et al.*, “CyberRAG: An Agentic RAG cyber attack classification and reporting tool,” *Future Generation Computer Systems*, vol. 176, p. 108186, 2026.
- [14] A. Abdennebi *et al.*, “LanG – A Governance-Aware Agentic AI Platform for Unified Security Operations,” *arXiv preprint arXiv:2604.05440*, 2026.
- [15] M. Analyst *et al.*, “Policy-Guided Threat Hunting: An LLM enabled Framework with Splunk SOC Triage,” *arXiv preprint arXiv:2603.23966*, 2026.