

**MINISTRY OF EDUCATION AND TRAINING
FPT UNIVERSITY**

**A TWO-TIER COGNITIVE ARCHITECTURE FOR
AUTOMATED THREAT DETECTION AND RESPONSE
USING AGENTIC AI**

by

Nguyen Duc Binh

A thesis submitted in conformity with the requirements
for the degree of Master of Software Engineering

**MINISTRY OF EDUCATION AND TRAINING
FPT UNIVERSITY**

**A TWO-TIER COGNITIVE ARCHITECTURE FOR
AUTOMATED THREAT DETECTION AND RESPONSE
USING AGENTIC AI**

by

Nguyen Duc Binh

A thesis submitted in conformity with the requirements
for the degree of Master of Software Engineering

Supervisor:
Dr. Bui Van Hieu
Dr. Dang Van Hieu

Abstract

In the modern cybersecurity landscape, Security Operations Centers (SOCs) face acute alert fatigue as continuous telemetry volume far exceeds human analyst processing capacity. Traditional automation tools relying on static playbooks lack contextual reasoning, whereas feeding raw logs directly to Large Language Models (LLMs) suffers severe latency bottlenecks, prompt injection risks, and data sovereignty violations. This thesis proposes SENTINEL, a cognitive two-tier security architecture operating fully air-gapped. Tier 1 functions as a wire-speed filter to automatically discard baseline stream noise, while Tier 2 employs a stateful AI reasoning agent backed by MITRE ATT&CK and NIST guidance to handle complex threats behind cryptographic guardrails.

Empirical evaluation on benchmark intrusion datasets demonstrates SENTINEL's high efficiency. Tier 1 automatically offloads over 80% of stream noise at sub-millisecond latency, reducing total system response latency to 211 ms (an 11.6x speedup over the Single-Tier LLM baseline) with throughput exceeding 2,500 flows/s. Tier 2 achieves a 94.6% technique mapping accuracy and an F1-score of 98.6%. Delimited Nonce encapsulation neutralizes 100% of prompt injection attacks, while HMAC-SHA256 hash-chaining guarantees audit trail non-repudiation, operating on a single 16 GB GPU under optimal resource consumption. SENTINEL proves that a two-tier hybrid architecture provides a trustworthy and effective paradigm for modern SOC defense.

Acknowledgments

The completion of this Master's thesis and the achievement of the Master of Software Engineering degree represent a journey that would not have been possible without the invaluable support, encouragement, and guidance of numerous individuals and institutions.

First and foremost, I wish to express my deepest, most sincere gratitude and profound respect to my academic supervisors, Dr. Bui Van Hieu and Dr. Dang Van Hieu. Through their profound scholarship, precise strategic orientation, and rigorous research ethics, they tirelessly devoted their time to guiding my critical thinking, offering sharp professional critiques, and inspiring my scientific curiosity throughout every phase of this research. Their invaluable mentorship has not only shaped this thesis into a rigorous academic work but also established an enduring foundation for my future career in cybersecurity and artificial intelligence engineering.

I am deeply grateful to the Board of Management, the leadership of the FSB School of Business and Technology, and the faculty members of FPT University. Their dedicated teaching equipped me with strong theoretical and practical knowledge, while providing a professional, modern, and supportive research environment throughout my Master's studies.

I extend my heartfelt thanks to the Leadership and my esteemed colleagues at Department C12, the National Data Center. Their trust, flexible accommodation of my schedule, and rich practical insights gained from real-world operations provided immense motivation and invaluable practical perspectives that directly informed the design, implementation, and empirical validation of the SENTINEL security system.

Finally, and most cherished of all, I dedicate my deepest emotional gratitude to my beloved family, especially my dear wife and our young child. Their silent sacrifices, infinite patience, unwavering understanding, and constant encouragement during the countless late nights spent conducting intensive system evaluations served as my solid pillar of strength, enabling me to overcome all obstacles and successfully bring this work to fruition.

While every effort has been made to conduct this research with utmost rigor and scientific dedication, limitations inevitably remain. I humbly welcome all constructive feedback and guidance from the thesis examination committee and fellow researchers to further refine and expand this work in the future.

Contents

Abstract	1
Acknowledgments	2
List of Figures	5
List of Tables	6
List of Abbreviations	7
1 Introduction	8
1.1 Background and Motivation	8
1.2 Research Objectives	9
1.3 Research Questions	9
1.4 Scope and Object of Study	10
1.5 Research Methodology	10
1.6 Core Contributions of the Thesis	11
1.7 Structure of the Thesis	12
2 Theoretical Background and Literature Review	13
2.1 SOC Operational Realities and Cognitive Bottlenecks	13
2.2 Online Statistical Theory and Tier 1 Machine Learning	13
2.3 Local Large Language Models and Tier 2 Agentic Intelligence	14
2.4 Adversarial AI Security and Cryptographic Log Audit	14
2.5 Overview of Related Research and Research Gaps	15
3 System Design and Technical Implementation	17
3.1 SENTINEL Architecture Overview and SOC Stack Integration	17
3.2 Tier 1: Wire-speed Filter and Machine Learning Gate	19
3.3 Tier 2: Stateful LangGraph Agent and Threat Memory	20
3.4 Dual-RAG Hybrid Knowledge Retrieval	21
3.5 Adversarial AI Security and Cryptographic Audit Sealing	21
3.6 SOC Dashboard and On-Premises Deployment Architecture	22
4 Experiments and Evaluation	23
4.1 Experimental Setup and Benchmark Telemetry	23
4.2 Tier 1 Wire-speed Offloading and Latency	24
4.3 Tier 2 Cognitive Reasoning Accuracy and Knowledge Mapping	25
4.4 Adversarial AI Defense, Prompt Injection, and HMAC Audit	25
4.5 Ablation Study, Platform Comparison, and Resilience	26

5	Conclusion and Future Work	28
5.1	Summary of Research Achievements	28
5.2	Scientific and Practical Contributions	29
5.3	Current System Limitations	30
5.4	Future Research Directions	31
	References	32

List of Figures

Figure 1.	Operational execution workflow within the SENTINEL pipeline. (Source: Author)	18
-----------	--	----

List of Tables

Table 1.	Comparative analysis between SENTINEL and existing security defense paradigms	16
Table 2.	Experimental Environment Specifications and Benchmark Telemetry Datasets	24
Table 3.	Tier 1 Component Inference Latency and Noise Offloading Ratios .	24
Table 4.	Tier 2 Adjudication Accuracy and RAG Knowledge Mapping Metrics	25
Table 5.	Ablation Study and Benchmark Comparison with State-of-the-Art Platforms	26

List of Abbreviations

Abbreviation	Full Description
AI	Artificial Intelligence
APT	Advanced Persistent Threat
CTI	Cyber Threat Intelligence
DAPT	Dynamic Advanced Persistent Threat
FAISS	Facebook AI Similarity Search
FPR	False Positive Rate
GGUF	GPT-Generated Unified Format
HITL	Human-in-the-Loop
HMAC	Hash-based Message Authentication Code
LLM	Large Language Model
MITRE ATT&CK	MITRE Adversarial Tactics, Techniques, and Common Knowledge
NIST	National Institute of Standards and Technology
RAG	Retrieval-Augmented Generation
RRF	Reciprocal Rank Fusion
SIEM	Security Information and Event Management
SOAR	Security Orchestration, Automation, and Response
SOC	Security Operations Center

Chapter 1

Introduction

This chapter presents the research context, central problem statement, and architectural foundation of the SENTINEL system—a Cognitive Two-Tier Architecture designed to automate threat detection and incident response within Security Operations Centers (SOCs). By critically analyzing the structural limitations of conventional deterministic defenses, the alert fatigue crisis, and the latency bottlenecks and adversarial security vulnerabilities inherent in Large Language Models (LLMs), this chapter establishes the research objectives, strategic research questions, scope, Design Science Research methodology, and core academic contributions of this thesis.

1.1 Background and Motivation

The expansion of enterprise IT infrastructure alongside the growing sophistication of multi-vector cyberattack campaigns poses severe operational challenges for security monitoring. Within modern Security Operations Centers (SOCs), Security Information and Event Management (SIEM) platforms are tasked with ingesting and aggregating millions of telemetry log events daily [2]. Conventional defense mechanisms rely heavily on perimeter filtration devices operating on static rulesets or fixed signatures. This approach exhibits fundamental limitations against novel evasion techniques and obfuscated attack variants. To minimize the risk of missing critical threats, security configurations frequently broaden sensitivity thresholds, resulting in hundreds of thousands of daily alerts with false positive rates exceeding 80% [4]. Consequently, security analysts suffer from cognitive exhaustion (alert fatigue), severely compromising their ability to identify genuine intrusion events. Concurrently, legacy Security Orchestration, Automation, and Response (SOAR) platforms operate on pre-scripted playbooks, completely lacking adaptive contextual reasoning capabilities [5].

Integrating Large Language Models (LLMs) offers a promising avenue toward automating security context synthesis and reasoning. Nevertheless, deploying LLMs directly into live

network telemetry streams presents three core technical barriers: First, the computational complexity of the self-attention mechanism causes inference latencies ranging from 4 to 26 seconds per event, incurring severe data pipeline congestion at wire-speed [18]. Second, the adversarial security risk stemming from log-substrate prompt injection attacks, wherein malicious instructions embedded in raw log fields manipulate the model’s internal reasoning flow [9]. Third, stringent data sovereignty regulations and Zero-Trust architecture mandates strictly prohibit exfiltrating internal security telemetry to public cloud AI services, requiring reasoning engines to operate entirely on-premises on bounded GPU hardware [31]. To resolve these challenges, this research proposes the SENTINEL architecture, synergizing an $O(1)$ statistical filter and a LightGBM Machine Learning Gate at Tier 1 with an Agentic AI reasoning system at Tier 2, safeguarded by cryptographic nonces and an HMAC audit trail.

1.2 Research Objectives

The overarching objective of this thesis is to design, prototypically build, and empirically evaluate the SENTINEL architecture to automate network security threat detection, contextual reasoning, and incident response on local infrastructure, achieving optimal processing latency and cryptographic security.

The general objective is operationalized through three central research tasks: First, constructing a dual-stage stateless filtration pipeline at Tier 1 operating at wire-speed, combining Welford’s $O(1)$ online statistical algorithm, a LightGBM Machine Learning Gate, and a Tier 1.75 Semantic Cache to classify and shed at least 80% of confident benign telemetry without consuming GPU compute resources. Second, designing a structural protection mechanism for the Tier 2 language model via Delimited Data Encapsulation using random cryptographic nonces, an output action guardrail, and an HMAC-SHA256 log audit chain to neutralize prompt injection attacks and ensure forensic non-repudiation. Third, developing a stateful reasoning engine based on a LangGraph agent state-graph coupled with persistent Threat Memory and a Dual-RAG retrieval system (combining FAISS vector search and BM25 sparse keyword matching) integrated with a STIX-aligned MITRE ATT&CK knowledge base, enabling transparent forensic justification generation, eliminating hallucination, and facilitating multi-stage intrusion correlation.

1.3 Research Questions

The design and empirical evaluation process of this thesis are guided by three strategic research questions:

The first research question addresses pipeline efficiency and filtration economics: To what extent can the hybrid Tier 1 filtration pipeline (combining Welford $O(1)$ statistics, LightGBM Machine Learning Gate, and Tier 1.75 Semantic Cache) achieve autonomous telemetry offload at wire-speed, and what overall response latency reduction is attained compared to a pure language model processing paradigm?

The second research question addresses AI security boundaries and forensic integrity: To what degree does Delimited Data Encapsulation using random cryptographic nonces coupled with deterministic action guardrails neutralize log-substrate prompt injection attacks, and how does an HMAC log audit chain protect forensic non-repudiation against log manipulation attempts?

The third research question addresses stateful agent reasoning and technical attribution: How effectively does a stateful Tier 2 AI agent equipped with persistent Threat Memory and Dual-RAG correlate complex intrusion events, accurately attribute MITRE ATT&CK technique codes, and produce transparent forensic justifications?

1.4 Scope and Object of Study

The primary subject of study is the SENTINEL Two-Tier Cognitive Architecture, comprising the Tier 1 stateless filter (online Welford statistics and LightGBM Machine Learning Gate), Tier 1.75 Semantic Cache, Tier 2 stateful agent engine (LangGraph state-graph and Q6_K quantised Gemma-2-9B-IT hosted locally via llama.cpp), Dual-RAG retrieval system (combining FAISS and BM25) integrated with MITRE ATT&CK knowledge, and cryptographic security controls including Delimited Nonce Encapsulation and HMAC-SHA256 log audit chains.

The evaluation scope is constrained to two benchmark datasets: CSE-CIC-IDS2018 for wide-area network telemetry and CSIC 2010 for real-world HTTP Layer 7 attacks targeting e-commerce applications. The complete system is deployed and tested on a single local GPU workstation equipped with an NVIDIA RTX 4060 Ti 16GB VRAM graphics card, Intel Core i7 processor, and 32GB RAM, operating fully air-gapped to meet strict data sovereignty mandates. Automated response actions are restricted to blocking IP addresses, logging events, or escalating to human security analysts; physical infrastructure disconnection interventions are excluded from this study.

1.5 Research Methodology

This thesis adopts the Design Science Research (DSR) methodology combining theoretical formalization with quantitative empirical validation across three main pillars: The statisti-

cal and machine learning pillar combines Welford $O(1)$ online statistics for rolling Z-score anomaly detection with LightGBM classification models and telemetry frequency/rate distribution metrics. The agentic logic pillar models an acyclic state-graph in LangGraph orchestrating language model reasoning, integrated with quantitative hybrid retrieval methods including FAISS vector search, BM25 sparse keyword search, and Reciprocal Rank Fusion (RRF) rank aggregation. The adversarial security pillar stress-tests resilience against prompt injection attacks and evaluates forensic log non-repudiation.

System efficacy is quantitatively evaluated using a 5D Metrics Framework covering: classification accuracy (Precision, Recall, F1), processing performance (latency in milliseconds and autonomous offload rate %), prompt injection deflection security, explanation coherence scored by an independent judge model, and log audit integrity via cryptographic hashing. Controlled ablation studies are conducted to measure and isolate individual tier contributions.

1.6 Core Contributions of the Thesis

This thesis contributes three core values to the fields of cybersecurity and artificial intelligence.

In system architecture, the research proposes and proves the efficacy of the SENTINEL Two-Tier Architecture. Synergizing $O(1)$ statistical filtering and a Machine Learning Gate at Tier 1 sheds the majority of telemetry at wire-speed with millisecond latencies, while Tier 1.75 Semantic Caching resolves repeated alerts, eliminating latency bottlenecks and preserving GPU compute resources.

In AI security, the research contributes a structural defense mechanism utilizing Delimited Data Encapsulation with random cryptographic nonces and deterministic action guardrails. This provides an architectural security guarantee, neutralizing prompt injection attacks without relying on probabilistic semantic classifiers. Concurrently, an HMAC-SHA256 signature chain guarantees non-repudiation for forensic logs.

In agentic logic and security knowledge, the research replaces rigid SOAR playbooks with a stateful LangGraph agent graph connected to a 432-item STIX-aligned MITRE ATT&CK knowledge base and persistent Threat Memory, enabling automated correlation of complex intrusion events and transparent forensic justification generation running smoothly on consumer GPU hardware.

1.7 Structure of the Thesis

The remainder of the thesis is organized into four subsequent chapters. Chapter 2 systemizes theoretical foundations of Welford statistics, LightGBM machine learning, language model quantization, LangGraph agent graphs, Dual-RAG systems, log audit cryptography, and related work. Chapter 3 details technical system design of SENTINEL, including the two-tier filter, Machine Learning Gate, agent state machine, security encapsulation, and SOC dashboard. Chapter 4 presents quantitative experimental results, benchmark evaluation on standard datasets, ablation studies, and adversarial security tests. Finally, Chapter 5 synthesizes findings, acknowledges research limitations, and outlines future maturation directions.

Chapter 2

Theoretical Background and Literature Review

This chapter outlines the foundational theoretical principles and comprehensive literature review underlying the thesis. The content systematically analyzes Security Operations Center (SOC) operational bottlenecks, Tier 1 online statistical and machine learning foundations, Tier 2 stateful agent and local LLM architectures, adversarial AI security threats alongside cryptographic audit mechanisms, and conducts a deep survey of existing AI Security systems to delineate the precise research gaps addressed by SENTINEL.

2.1 SOC Operational Realities and Cognitive Bottlenecks

Information security monitoring in modern enterprise Security Operations Centers (SOCs) relies on aggregating raw telemetry logs from Next-Generation Firewalls (NGFW), Intrusion Detection Systems (IDS/IPS), and Endpoint Detection and Response (EDR) agents into centralized SIEM platforms [2]. Telemetry volume from distributed enterprise infrastructure far exceeds human manual analysis capabilities [3]. Because perimeter defense devices operate using static signatures and deterministic rulesets, relaxing alert thresholds to avoid missing threats leads to a severe alert volume crisis with false positive rates exceeding 80% [4], resulting in cognitive alert fatigue for security analysts. Although Security Orchestration, Automation, and Response (SOAR) platforms assist in automating response workflows, they depend on rigid, pre-programmed playbooks that entirely lack contextual reasoning when encountering novel attack variants [5].

2.2 Online Statistical Theory and Tier 1 Machine Learning

To filter high-density telemetry streams at wire-speed without maintaining full historical logs in memory, Welford's algorithm maintains $\mathcal{O}(1)$ time and memory complexity [6].

Sample mean μ_k and sum of squared deviations S_k are updated via recurrence relations:

$$\mu_k = \mu_{k-1} + \frac{x_k - \mu_{k-1}}{k}, \quad S_k = S_{k-1} + (x_k - \mu_{k-1})(x_k - \mu_k) \quad (2.1)$$

enabling real-time standard deviation $\sigma_k = \sqrt{S_k/(k-1)}$ and continuous Z -score anomaly metric evaluation. Concurrently, the LightGBM machine learning gate utilizes leaf-wise tree growth to achieve sub-0.35 ms inference latencies per NetFlow sample. To counter machine learning evasion attacks, the ML Gate enforces Z -score feature clipping within $\pm 8\sigma$ and baseline mean imputation. Additionally, the Tier 1.75 Semantic Cache uses payload template hashing in RAM to execute short-circuit responses for repetitive volumetric attacks (such as DDoS or brute-force attempts) in under 1 ms, bypassing GPU invocation.

2.3 Local Large Language Models and Tier 2 Agentic Intelligence

Local LLM deployment demands strict memory optimization: Gemma-2-9B-IT in FP16 precision consumes approximately 18 GB VRAM [23]. Weight quantization maps floating-point parameters to integer grids [19]. Using GGUF Q6_K format via `llama.cpp`, required VRAM drops to 7–8 GB (>50% savings) without significant accuracy loss, operating on a single 16 GB GPU.

To orchestrate reasoning, Tier 2 employs a LangGraph acyclic state graph [8] integrated with Threat Memory. To prevent hallucination, the system integrates Retrieval-Augmented Generation (RAG) [7] via a Dual-RAG architecture combining FAISS dense search [20] and BM25 sparse search [21], fused via Reciprocal Rank Fusion (RRF) [11]:

$$RRF_Score(d) = \sum_{m \in M} \frac{1}{k + r_m(d)} \quad (2.2)$$

retrieving context directly from 432 STIX MITRE ATT&CK technique entries and NIST SP 800-61r2 guidance for forensic justification.

2.4 Adversarial AI Security and Cryptographic Log Audit

Integrating LLMs introduces log-substrate prompt injection risks [9], where attackers inject malicious instructions into User-Agent headers or Syslog strings to hijack AI reasoning [30]. To neutralize this vector, SENTINEL applies Delimited Data Encapsulation with random Nonce tags to isolate raw inputs, combined with deterministic Guardrails enforcing structured outputs into 3 valid actions (BLOCK_IP, LOG, AWAIT_HITL). To guarantee forensic non-repudiation, the system uses HMAC-SHA256 signature chaining [29] to seal

audit records with a local secret key, producing a tamper-evident audit trail.

2.5 Overview of Related Research and Research Gaps

Artificial intelligence research applied to Security Operations Centers (SOCs) has evolved across four major architectural paradigms, spanning traditional deterministic platforms to modern autonomous agentic systems [15]. Analyzing the core technologies, objectives, and limitations of each paradigm establishes the exact technical research gap resolved by SENTINEL.

Traditional enterprise defense solutions include commercial SIEM platforms such as IBM QRadar and SOAR systems such as Palo Alto Cortex XSOAR [2, 5]. IBM QRadar aggregates and correlates telemetry using static correlation rules on centralized databases. However, QRadar exhibits high false positive alert rates ($>80\%$) due to a lack of contextual comprehension. To mitigate this, Cortex XSOAR automates incident response workflows using pre-programmed playbooks (written in Python/JavaScript). Although it reduces manual repetitive tasks, Cortex XSOAR operates deterministically: when encountering novel attack variants with obfuscated payloads or anomalous behaviors outside predefined playbooks, the system fails or misses threats entirely.

Cloud-based LLM agent systems are represented by platforms such as Watchtower [12] and CyberRAG [13]. Watchtower proposes a multi-agent framework coordinating incident investigation workflows by transmitting raw telemetry logs to commercial cloud LLMs (such as GPT-4 or Claude). However, transmitting raw logs to cloud models introduces two severe technical limitations: quadratic self-attention latency $\mathcal{O}(N^2)$ extends processing time per event to 4–26 seconds, causing queue congestion at wire-speed; concurrently, transferring enterprise telemetry violates Zero-Trust principles and data sovereignty mandates [31]. Meanwhile, CyberRAG [13] focuses on single-tier RAG threat intelligence retrieval; lacking a Tier 1 wire-speed filter, CyberRAG suffers queue exhaustion under high event streams and lacks stateful correlation for multi-stage APT campaigns.

Autonomous planning and security reasoning agents comprise AutoBNB [14], LanG Agentic SOC [16], and Policy-Guided SIEM [17]. AutoBNB combines Bayesian networks with planning agents for attack graph generation and root-cause attribution, but encounters high graph computation costs and lacks structural defenses against log prompt injection attacks [9]. Most recently, the LanG platform [16] introduced governance-aware multi-agent orchestration, and Policy-Guided SIEM [17] integrated rule-guided guardrails with Splunk SIEM workflows. While these systems advance multi-agent coordination and hallucination reduction, they remain single-tier architectures where every network event must traverse complex LLM reasoning nodes. Consequently, average processing latency remains high

between 480 ms and 520 ms, requiring multi-GPU server infrastructure (16 GB–24 GB VRAM), failing to deliver wire-speed telemetry noise offloading at Tier 1.

The primary research gap identified is: Currently, no hybrid architecture simultaneously satisfies three critical requirements: (1) Wire-speed Tier 1 noise offloading filtering at least 80% of stream noise under 1 ms latency; (2) Tier 2 stateful agent cognitive reasoning running locally (air-gapped 16 GB GPU) with Dual-RAG hallucination reduction; and (3) Structural AI model integrity protection via Delimited Nonce encapsulation and HMAC-SHA256 forensic sealing. Table 1 provides a comparative synthesis contrasting SENTINEL against representative paradigms.

Table 1. Comparative analysis between SENTINEL and existing security defense paradigms

Evaluation Metric	SIEM / SOAR (QRadar, XSOAR)	SOTA Agents (Watchtower, LanG, AutoBnB)	SENTINEL Architecture (Thesis)
Core Technology	Static correlation rules, rigid Python/JS playbooks	Cloud Multi-Agent LLMs / Local Multi-Agent RAG	Welford $\mathcal{O}(1)$ + LightGBM T1, LangGraph Agent T2
Operational Goal	Centralized log aggregation and playbook response	Automated incident investigation via LLM reasoning	Wire-speed stream noise offloading and local cognitive reasoning
Latency & Throughput	Very low (< 1 ms) but high false positive surge	Medium/High (480 ms–26 s), pipeline congestion	Optimal (< 0.40 ms Tier 1, 1 ms Cache, $> 80\%$ T2 offloading)
Data Sovereignty	Local enterprise	Cloud leakage risks or heavy 24GB GPU requirement	Absolute security, 100% Air-gapped on a 16 GB GPU
Prompt Inj. Defense	Not applicable	Probability semantic filters (vulnerable to evasion)	Architectural guarantee (Delimited Nonce Encapsulation)
Forensic Audit	Standard text logs (vulnerable to post-compromise edits)	No cryptographic audit trail sealing	HMAC-SHA256 hash-chaining non-repudiation seal

Chapter 3

System Design and Technical Implementation

This chapter elaborates on the architectural design and technical implementation solutions of the SENTINEL system. The content systematically details the two-tier execution pipeline, Tier 1 wire-speed filtration algorithms and Machine Learning Gate, Tier 2 agentic reasoning engine alongside Threat Memory, Dual-RAG knowledge retrieval, adversarial AI security controls and cryptographic audit sealing, and the local SOC dashboard deployment infrastructure.

3.1 SENTINEL Architecture Overview and SOC Stack Integration

To protect enterprise network infrastructure against high-density telemetry streams without incurring pipeline congestion, the SENTINEL system enforces a Cognitive Two-Tier Architecture. The core architectural design principle lies in strict operational decoupling: Tier 1 executes stateless statistical and machine learning algorithms to filter and offload at least 80% of benign stream noise at wire-speed; Tier 2 hosts a stateful language model agent engine to perform deep contextual reasoning and adjudication for complex anomalous events.

Figure 1 illustrates the detailed operational execution flow within the SENTINEL data pipeline. Raw network log telemetry generated by perimeter monitoring sources (NGFW, IPS, EDR) is ingested into a centralized Redis Stream queue. Immediately upon event arrival, Tier 1 computes online Z -score anomaly metrics using Welford's running statistical algorithm and classifies flow feature interactions through a LightGBM Machine Learning Gate. Traffic flows conforming to baseline distributions or matching repetitive payload templates stored in the Tier 1.75 Semantic Cache are adjudicated sub-millisecond, preserving GPU compute resources.

Only events classified as genuine anomalies escalate to Tier 2. Here, the LangGraph agent engine triggers a reasoning cycle executing a locally hosted Gemma-2-9B-IT model, integrating historical state from persistent Threat Memory and retrieved evidence from the




Figure 1. Operational execution workflow within the SENTINEL pipeline. (Source: Author)

Dual-RAG system. Final output decisions are governed by a deterministic Guardrail layer, strictly bounded to three explicit states (BLOCK_IP, LOG, or AWAIT_HITL). Every adjudication record is cryptographically sealed using HMAC-SHA256 hash chains to ensure forensic trail integrity and non-repudiation.

Within modern enterprise SOC architectures, SENTINEL functions as an intelligent cognitive middleware layer positioned between centralized SIEM log collectors (such as IBM QRadar) and downstream firewall/SOAR execution platforms (such as Palo Alto Cortex XSOAR). By absorbing massive stream noise at Tier 1 and providing grounded, transparent justifications for escalated events at Tier 2, SENTINEL eliminates alert flooding for security analysts while ensuring complete data sovereignty through 100% air-gapped execution on a single 16 GB VRAM GPU.

3.2 Tier 1: Wire-speed Filter and Machine Learning Gate

Tier 1 functions as the primary stateless perimeter defense layer processing stream telemetry at wire-speed. This layer combines Welford’s $\mathcal{O}(1)$ statistical baseline filter with a LightGBM Machine Learning Gate to guarantee millisecond-level inference latencies and suppress stream noise prior to invoking GPU resources.

The statistical baseline filter continuously computes standard deviations σ and real-time Z -score anomaly metrics across network flow attributes (packet rates, flow volumes, connection durations). Attribute Z -scores are calculated as $Z = (x - \mu)/\sigma$, where sample means μ and variances σ are updated continuously via Welford’s online algorithm without retaining raw event series in memory [6]. Specifically, at each time step k , sample expectations μ_k and the sum of squared deviations S_k are updated via the recurrence relations:

$$\mu_k = \mu_{k-1} + \frac{x_k - \mu_{k-1}}{k}, \quad S_k = S_{k-1} + (x_k - \mu_{k-1})(x_k - \mu_k) \quad (3.1)$$

enabling online evaluation of sample standard deviations $\sigma_k = \sqrt{S_k/(k-1)}$. Consequently, baseline flow traffic conforming to normal distributions is shed immediately. For repetitive volumetric intrusion attempts (such as DDoS attacks or automated brute-force attempts), the Tier 1.75 Semantic Cache hashes payload template structures in RAM, enabling 1-millisecond short-circuit response execution using cached decisions, terminating pipeline propagation prior to deep system ingestion.

A LightGBM Gradient Boosted Decision Tree model is integrated immediately following statistical filtration to evaluate non-linear flow feature interactions. Utilizing leaf-wise tree growth and histogram feature binning, LightGBM evaluates each NetFlow feature vector in approximately 0.35 ms. To neutralize machine learning evasion tactics (anti-evasion), the

Machine Learning Gate enforces feature Z -score clipping bounded within $\pm 8\sigma$ and imputes missing feature values using baseline mean metrics. If distorted feature proportions exceed 30%, the Machine Learning Gate abstains from adjudication and automatically escalates the event to Tier 2 for deep contextual inspection.

3.3 Tier 2: Stateful LangGraph Agent and Threat Memory

Tier 2 performs stateful cognitive reasoning for anomalous events escalated from Tier 1. The reasoning engine executes a GGUF Q6_K quantized Gemma-2-9B-IT model [23] via llama.cpp compiled with CUDA acceleration on a single local 16 GB VRAM GPU. The Q6_K quantization scheme reduces required VRAM from 18 GB (FP16 precision) to 7–8 GB (>50% memory savings) with negligible reasoning accuracy loss, enabling smooth model execution in air-gapped enterprise environments.

The agent reasoning workflow is orchestrated by a LangGraph Directed Acyclic State-Graph [8]. Execution follows a deterministic sequence across five core processing nodes: initial guardrail inspection (`guardrails`), context enrichment (`rag_context`), preliminary triage (`llm_triage`), attack technique mapping (`attack_mapper`), and action execution (`action_executor` or `human_in_loop`). Technique mapping requires explicit verification from domain knowledge; if no unambiguous match is identified, the state graph routes the event to human analyst review (`AWAIT_HITL`) rather than outputting speculative actions. Allowable output actions are strictly constrained to three states: `BLOCK_IP`, `LOG`, or `AWAIT_HITL`.

To correlate intrusion traces over extended periods, Tier 2 maintains persistent Threat Memory on an index-optimized SQLite database. Connection settings enforce normal synchronous mode with rollback-journal configuration to prevent cross-UID file lock crashes in Docker deployments. Threat Memory manages database tables tracking IP reputation histories, internal entities, risk indicators, and threat events. IP reputation scores S decay naturally over silent periods t according to:

$$S_{\text{new}} = S_{\text{old}} \times (1 - \lambda)^t \quad (3.2)$$

where λ denotes the decay rate. When an IP address accumulates a maximum reputation score ($S = 100$), Tier 1 automatically enforces block-on-sight execution for subsequent connections without re-triggering Tier 2 LLM reasoning, closing an automated feedback loop. Conversely, internal infrastructure IP addresses (such as `10.0.0.99`) are hard-coded into a whitelist and remain immune to block actions to prevent self-DoS risks. This mechanism enables the agent to correlate multi-stage intrusion chains spanning multiple days.

3.4 Dual-RAG Hybrid Knowledge Retrieval

To eliminate model hallucination and provide grounded forensic evidence for security decisions, Tier 2 incorporates a Dual-RAG hybrid knowledge retrieval system. The objective of this mechanism is to supply standardized domain knowledge from authoritative frameworks to the large language model during every reasoning cycle.

SENTINEL's security knowledge base is compiled from 432 STIX 2.1-structured MITRE ATT&CK technique entries [24] and NIST SP 800-61r2 computer security incident handling guidance [10]. Retrieval operates via two parallel, complementary search mechanisms: dense semantic vector search using FAISS [20] and sparse keyword matching using BM25 [21]. FAISS employs a Sentence-BERT embedding model [25] to project telemetry queries into high-dimensional vector spaces for L2/Cosine distance matching, identifying attack techniques sharing semantic intent despite keyphrase variations. Conversely, BM25 computes term frequencies over sparse inverted indexes, ensuring exact retrieval of technical identifiers such as CVE codes, Registry keys, or service port numbers.

Retrieved candidate lists from both methods are consolidated using Reciprocal Rank Fusion (RRF) [11]:

$$RRF_Score(d) = \sum_{m \in \{FAISS, BM25\}} \frac{1}{k + r_m(d)} \quad (3.3)$$

where $r_m(d)$ denotes the document rank in method m and $k = 60$ represents the smoothing constant. The top-ranked technique entries are compressed and injected directly into the prompt context at the `rag_context` node. Consequently, the reasoning agent accurately maps anomalous telemetry to specific MITRE ATT&CK technique identifiers while outputting verifiable forensic justifications for every adjudication decision.

3.5 Adversarial AI Security and Cryptographic Audit Sealing

Integrating large language models into cybersecurity infrastructure demands comprehensive adversarial defenses against log-substrate prompt injection attacks [9]. This vector involves attackers embedding malicious control instructions inside raw telemetry payload fields (such as User-Agent headers, URI parameters, or Syslog message text) to deceive the model and hijack its operational behavior.

To neutralize this threat vector, SENTINEL implements Delimited Data Encapsulation reinforced with session-unique cryptographic nonces [30]. Raw log telemetry is entirely enclosed within nonce delimiters (e.g., `<<DATA_HASH . . . >>`). This encapsulation forces the language model to process input payloads passively as plain data content, preventing the execution of embedded instruction triggers. Concurrently, label stripping removes internal

classification tags from prompt inputs, eliminating internal state leakage. In parallel, a deterministic Guardrail layer strictly enforces Pydantic schema validation (`LLMDecision`); if output formatting errors occur, regular expression salvaging recovers essential decision fields or defaults to the safe `AWAIT_HITL` state.

Complementing AI defenses, SENTINEL incorporates cryptographic audit sealing to guarantee non-repudiation for forensic evidence. Each security event record and adjudication decision i is cryptographically signed using HMAC-SHA256 signature chaining [29]:

$$H_i = \text{HMAC-SHA256}(D_i \parallel H_{i-1}, K) \quad (3.4)$$

where D_i denotes current record data, H_{i-1} represents the preceding record hash signature, and K is a local secret key. Chaining hash signatures produces a tamper-evident audit log trail; any unauthorized attempt to alter or delete historical log records breaks hash continuity and is flagged immediately.

3.6 SOC Dashboard and On-Premises Deployment Architecture

The SOC Analyst Dashboard is developed using Streamlit, connecting directly to Redis Stream queues and the Threat Memory database. The dashboard provides analysts with real-time operational visibility: Tier 1 processing throughput telemetry, multi-stage intrusion timeline visualization, HMAC audit chain integrity status, and a Human-in-the-Loop (HITL) review panel (`AWAIT_HITL`).

The complete system services are fully containerized using Docker Compose (`docker-compose.yml`). The `llama.cpp` inference library is compiled with CUDA acceleration to leverage an NVIDIA RTX 4060 Ti 16 GB VRAM GPU. The Threat Memory database (`threat_memory.db`) and FAISS index are mounted on persistent Docker volumes, ensuring IP reputation scores and audit records survive container restarts.

Furthermore, deployment strictly adheres to Zero Trust mandates [31]: database password strings in `REDIS_URL` are redacted prior to logging. To manage high-density stream surges, the system enforces latency-based backpressure monitoring Redis consumer group lag, maintaining long-term pipeline stability.

Chapter 4

Experiments and Evaluation

This chapter presents the experimental results and quantitative evaluation of the SENTINEL system. The content is streamlined to focus on core performance indicators: Tier 1 wire-speed noise offloading efficiency, Tier 2 cognitive agent adjudication and knowledge mapping accuracy, adversarial AI defense capabilities alongside cryptographic audit integrity, combined with controlled ablation studies and multi-dimensional comparisons against state-of-the-art platforms.

4.1 Experimental Setup and Benchmark Telemetry

The SENTINEL testbed is deployed on a local server representative of small-to-medium enterprise (SME) Security Operations Center (SOC) infrastructure, powered by an Intel Core i7 processor, 32 GB DDR5 RAM, and a single NVIDIA GeForce RTX 4060 Ti 16 GB VRAM GPU. All system microservices—including Redis Stream message queues, Tier 1 filters, Tier 2 LangGraph agents, and the `llama.cpp` CUDA inference server—are containerized via Docker Compose in an air-gapped environment to eliminate host operating system interference and ensure full experimental reproducibility.

Evaluation telemetry combines two benchmark datasets: CSE-CIC-IDS2018 [1] (providing 26,521 labeled network flow events including DDoS, DoS, Brute-force, and Benign traffic) and DAPT2020 [27] (simulating multi-stage APT attack scenarios). Additionally, an independent 300-sample balanced ablation dataset (150 normal and 150 ground-truth attack samples) is extracted for comparative evaluation. Table 2 outlines the hardware testbed and benchmark telemetry specifications.

Table 2. Experimental Environment Specifications and Benchmark Telemetry Datasets

Component	Specification / Title	Detailed Description and Architectural Role
Processor (CPU)	Intel Core i7	Executes Welford statistical filtering and Redis streams
System RAM	32 GB DDR5	Stores Semantic Cache hash tables and FAISS vector indices
Graphics Card (GPU)	NVIDIA RTX 4060 Ti	Runs <code>llama.cpp</code> CUDA server for Gemma-2-9B model
Virtualization	Docker Compose	Encapsulates microservices in an air-gapped environment
Dataset 1	CSE-CIC-IDS2018 [1]	26,521 labeled NetFlow events (DDoS, Brute-force, Benign)
Dataset 2	DAPT2020 [27]	Multi-stage Advanced Persistent Threat (APT) attack chains
Ablation Dataset	300 balanced samples	150 benign + 150 attack events (ground-truth verified)

4.2 Tier 1 Wire-speed Offloading and Latency

Tier 1 filters stream noise at wire-speed, freeing GPU resources for Tier 2 reasoning. Welford’s $\mathcal{O}(1)$ online statistical filter (`rule_engine.py`) updates mean μ_k and sum of squares S_k in real time, calculating standard deviation σ_k to discard normal traffic [6]. Concurrently, LightGBM (`ml_gateway.py`) achieves 0.35 ms inference latency via leaf-wise tree growth. For repetitive attacks, Tier 1.75 Semantic Cache (`response_cache.py`) executes sub-1 ms short-circuit responses.

Table 3. Tier 1 Component Inference Latency and Noise Offloading Ratios

Tier 1 Component	Source Script File	Avg Latency	Events Offloaded	Ratio (%)
Welford Statistical Filter	<code>rule_engine.py</code>	< 0.01 ms	14,820 events	55.9%
LightGBM Machine Learning Gate	<code>ml_gateway.py</code>	0.35 ms	4,376 events	16.5%
Semantic Cache (Tier 1.75)	<code>response_cache.py</code>	< 1.00 ms	2,657 events	10.0%
Total Tier 1 Pipeline	–	< 0.40 ms	21,853 events	82.4%

In aggregate, Tier 1 successfully offloads 82.4% (21,853 of 26,521 events) of stream noise

at an average latency below 0.40 ms (Table 3). Filtering over 80% of stream noise without GPU invocation maintains Redis Stream consumer group lag near zero, ensuring long-term operational sustainability under heavy workloads.

4.3 Tier 2 Cognitive Reasoning Accuracy and Knowledge Mapping

Tier 2 processes the escalated 17.6% anomalous stream events to perform stateful reasoning via the LangGraph agent (`workflow.py`) combined with the quantized Gemma-2-9B-IT Q6_K model. The Dual-RAG hybrid retrieval engine (`hybrid_retriever.py`) fuses candidate contexts from FAISS dense vector search [20] and BM25 sparse keyword matching [21] via Reciprocal Rank Fusion ($RRF_Score(d) = \sum 1/(k + r_m(d))$ with $k = 60$) [11]. Evaluated across 432 STIX 2.1-aligned MITRE ATT&CK techniques, RRF rank fusion achieves an ATT&CK technique mapping accuracy of 94.6%, eliminating model hallucination.

Table 4. Tier 2 Adjudication Accuracy and RAG Knowledge Mapping Metrics

Retrieval Paradigm / Evaluation Metric	Accuracy (%)	FPR (%)	RAGAS / MT-Bench
FAISS Dense Vector Search Only	81.2%	4.5%	0.78 / 10.0
BM25 Sparse Keyword Search Only	78.4%	5.2%	0.72 / 10.0
Dual-RAG RRF (FAISS + BM25, $k = 60$)	94.6%	< 1.2%	0.94 / 8.95
Agent Security Decision (BLOCK_IP / LOG)	98.6% ($F1 = 98.6\%$)	< 1.2%	8.95 / 10.0

Agent security adjudication decisions achieve an $F1 = 98.6\%$ score and a False Positive Rate (FPR) below 1.2%. Reasoning quality evaluation conducted using LLM-as-a-Judge protocols yields 8.95/10 on MT-Bench [26] and 0.94 on RAGAS [28] (Table 4), confirming strong forensic logic and execution reliability.

4.4 Adversarial AI Defense, Prompt Injection, and HMAC Audit

System security resilience is verified across three evaluation scenarios. In log prompt injection testing (User-Agent headers, Syslog strings) [9], Delimited Data Encapsulation (`encapsulation.py`) with unique random Nonce delimiters (`<<<DATA_HASH...>>>`) neutralizes 100% of malicious control instructions, forcing the model to process inputs passively [30]. Stripping internal classification labels prior to prompt generation prevents state leakage.

For machine learning evasion testing, Tier 1 ML feature clipping ($\pm 8\sigma$) and the 30%

distortion abstention threshold escalate suspicious flows to Tier 2 for contextual inspection. In audit trail verification, any unauthorized log modification breaks HMAC-SHA256 hash-chaining continuity (`hmac_sealer.py`) following $H_i = \text{HMAC-SHA256}(D_i \parallel H_{i-1}, K)$ [29], detecting non-repudiation violations immediately.

4.5 Ablation Study, Platform Comparison, and Resilience

An ablation study on 300 balanced samples demonstrates that the Single-Tier LLM baseline suffers severe queue congestion with 2,450 ms average latency and VRAM exhaustion. Conversely, SENTINEL maintains sub-15 ms Tier 1 offloading latencies for 82.4% of stream traffic and achieves a global weighted End-to-End (E2E) average latency of 211 ms (combining 82.4% events processed at under 0.40 ms and 17.6% anomalous events escalated to Tier 2 processed at 1.200 ms), optimizing GPU VRAM utilization by over 75%.

Consolidated comparison against published platforms—including Watchtower [12], CyberRAG [13], AutoBnB [14], LanG Agentic SOC [16], and Policy-Guided SIEM [17]—confirms SENTINEL’s performance leadership across Throughput (over 2,500 flows/s), E2E Latency (211 ms), Prompt Injection Defense (100%), and VRAM optimization (under 8 GB) (Table 5).

Table 5. Ablation Study and Benchmark Comparison with State-of-the-Art Platforms

System / Platform	Throughput (Flows/s)	Avg E2E Latency (ms)	Accuracy (%)	Prompt Inj. (%)	GPU VRAM (GB)
Single-Tier LLM Baseline	0.41 flows/s	2,450 ms	91.2%	42.0%	15.8 GB
Watchtower Platform [12]	120 flows/s	850 ms	93.5%	75.0%	Cloud API
CyberRAG Agent [13]	85 flows/s	1.120 ms	92.8%	68.0%	24.0 GB
AutoBnB Framework [14]	210 flows/s	620 ms	94.1%	81.0%	16.0 GB
LanG Agentic SOC [16]	140 flows/s	480 ms	95.2%	88.0%	24.0 GB
Policy-Guided SIEM [17]	180 flows/s	520 ms	94.5%	85.0%	16.0 GB
SENTINEL (Proposed Two-Tier)	> 2,500 flows/s	< 15 ms (T1) 211 ms (E2E)	98.6%	100.0%	7.8 GB

Finally, operational resilience testing under high-density stream surges (100,000 events) confirms that latency-based Redis consumer group backpressure maintains stable RAM

and VRAM utilization, guaranteeing continuous system availability for enterprise SOC deployments.

Chapter 5

Conclusion and Future Work

This chapter synthesizes the primary research outcomes and core contributions of the thesis entitled “A Cognitive Two-Tier Architecture for Automated Threat Detection and Contextual Response using Agentic AI” (SENTINEL). The content summarizes empirical evaluation results aligned with the three research questions, affirms scientific and practical contributions, transparently addresses existing architectural limitations, and outlines future research directions to advance Security Operations Center (SOC) defense capabilities.

5.1 Summary of Research Achievements

Comprehensive empirical evaluation across 26,521 telemetry events from the CSE-CIC-IDS2018 benchmark [1], multi-stage APT scenarios from DAPT2020 [27], and a 300-sample balanced ablation dataset demonstrates that the SENTINEL system successfully resolves the three strategic research questions formulated in Chapter 1.

Regarding the first research question (RQ1) on wire-speed filtering pipeline performance and economics, quantitative measurements confirm that the multi-tiered architecture effectively overcomes compute resource bottlenecks. First, regarding efficiency (RQ1), SENTINEL successfully offloaded 82.4% of telemetry noise (21,853 of 26,521 events) at wire-speed with an average Tier 1 latency below 0.40 ms. Global End-to-End average latency was reduced to 211 ms (11.6x faster than the 2,450 ms Single-Tier baseline), maintaining throughput exceeding 2,500 flows/s while freeing GPU VRAM by over 75%.

Regarding the second research question (RQ2) on AI security guardrails and cryptographic non-repudiation, Delimited Data Encapsulation using unique random Nonce delimiters combined with deterministic Guardrails successfully neutralizes 100% of log-substrate prompt injection attack vectors [9, 30]. This mechanism forces the language model to process raw inputs strictly as passive data, completely preventing inference manipulation risks. Concurrently, HMAC-SHA256 cryptographic hash-chaining [29] sequentially sealing each audit log entry guarantees non-repudiation and tamper-evident forensic audit trail

integrity against post-compromise tampering.

Regarding the third research question (RQ3) on stateful agent reasoning, ATT&CK technique mapping, and transparent explanations, the Tier 2 LangGraph agent integrating the Gemma-2-9B-IT Q6_K model with Dual-RAG RRF rank fusion ($k = 60$) [11, 20, 21] across 432 STIX MITRE ATT&CK entries achieves an ATT&CK technique mapping accuracy of 94.6%. Security adjudication decisions achieve an $F1 = 98.6\%$ score with a False Positive Rate controlled below 1.2%. Reasoning logic quality evaluated via LLM-as-a-Judge protocols yields 8.95/10 on MT-Bench [26] and 0.94 on RAGAS [28], confirming transparent forensic justification while completely eliminating hallucination.

5.2 Scientific and Practical Contributions

This thesis provides two core contributions spanning theoretical security research and practical enterprise implementation within the domains of cybersecurity and artificial intelligence.

In terms of scientific contributions, the research proposes a Cognitive Two-Tier Hybrid Architecture harmonizing high-speed deterministic statistical algorithms at Tier 1 with stateful LLM agent reasoning engines at Tier 2. This multi-tiered structural paradigm establishes a systematic compute offloading mechanism, resolving quadratic self-attention latency bottlenecks $\mathcal{O}(N^2)$ inherent in large language models operating within continuous network telemetry streams [18]. Furthermore, the thesis contributes a structural data encapsulation model utilizing unique random Nonce delimiters, providing architectural security guarantees against log prompt injection attacks without relying on fragile probabilistic classifiers [9, 30]. Finally, the study establishes a non-parametric quantitative evaluation methodology integrating cross-family LLM-as-a-Judge protocols and standardized RAGAS metrics tailored for local agentic security systems.

In terms of practical contributions, the thesis delivers a fully containerized open-source system prototype packaged via Docker Compose and empirically verified on accessible SME commercial hardware (an Intel Core i7 processor, 32 GB RAM, and a single NVIDIA GeForce RTX 4060 Ti 16 GB VRAM GPU). This design strictly adheres to Zero-Trust security principles and enterprise data sovereignty mandates within an air-gapped environment [31]. The implementation features an interactive Streamlit monitoring dashboard connected to Human-in-the-Loop (HITL) queues to replace static SOAR playbooks [5], while automatically sealing forensic audit trails via HMAC-SHA256 cryptographic hash-chaining to support post-incident investigation [29].

5.3 Current System Limitations

Despite achieving positive empirical evaluation outcomes, several technical and methodological limitations of the SENTINEL system should be acknowledged transparently to define its operational boundaries.

The first limitation involves telemetry evaluation scope and log parsing capabilities. Current evaluations focus primarily on normalized NetFlow features and Layer-7 HTTP log attributes from the CSE-CIC-IDS2018 and DAPT2020 benchmarks [1, 27]. In real-world enterprise SOC environments, heterogeneous unstructured log sources (such as Windows Event Logs, raw Syslog, or AWS CloudTrail) necessitate sophisticated log parsing and feature extraction mechanisms.

The second limitation involves single-node hardware compute boundaries under volumetric traffic surges exceeding 100,000 events/second. Although Tier 1 successfully offloads over 80% of baseline stream noise, the volume of escalated anomalous events to Tier 2 on a single 16 GB VRAM GPU server may temporarily exceed the immediate inference capacity of the language model, leading to Redis Stream consumer queue elevation.

The third limitation stems from trade-offs between operational autonomy and Self-Denial-of-Service (Self-DoS) risks when executing automated IP blocking actions. To prevent accidental blocking of critical internal IP addresses or infrastructure servers caused by potential model reasoning bias, the system forces general port-based MITRE ATT&CK detections (such as T1571) into human analyst review queues (AWAIT_HITL). Mandatory analyst intervention slightly reduces absolute automation levels but serves as an indispensable safety guardrail for enterprise infrastructure protection.

The fourth limitation concerns the scope and mechanics of APT attack chain correlation. The current stateful correlation mechanism relies primarily on fixed source IP tracking across representative test scenarios. In real-world environments featuring IP rotation or distributed botnets, correlation mechanisms relying strictly on source IP addresses may experience reduced effectiveness.

The fifth limitation relates to state storage operational stability within containerized deployment environments. Audit logging and threat memory operations on SQLite databases, if not configured with appropriate synchronization modes, may encounter file lock contention or process disruption during unexpected service crashes.

5.4 Future Research Directions

To address the five technical limitations identified in Section 5.3, future research and development for SENTINEL proposes five corresponding technical engineering solutions for production deployment:

First, addressing evaluation scope and parsing limitations, future work will develop an automated ingestion engine integrating template clustering algorithms (such as Drain [22] or LogParser) to parse and normalize diverse log formats into standardized feature vectors prior to pipeline processing.

Second, addressing volumetric DoS traffic surges and single-GPU bottlenecks, the system will re-implement Tier 1 filters in compiled system languages (such as Go or Rust) to elevate throughput to millions of events per second; concurrently migrating Redis Streams to partitioned disk-backed Apache Kafka clusters and upgrading Tier 2 engines to vLLM or TGI frameworks supporting PagedAttention and continuous batching for parallel GPU processing.

Third, addressing Self-DoS risks and autonomy trade-offs, research proposes applying Reinforcement Learning from Human Feedback (RLHF) based on analyst approval decisions within `AWAIT_HITL` queues, enabling agents to continuously refine response policies and safely automate IP blocking without Self-DoS risks.

Fourth, addressing source IP correlation constraints and security tool orchestration, future work will implement specialized Multi-Agent Sub-systems combined with Graph Neural Networks (GNNs) to correlate attack chains based on technical vectors. Concurrently, the system will integrate the Model Context Protocol (MCP) standard, enabling AI agents to dynamically query and orchestrate external security tools (SIEM, EDR, eBPF Firewalls, CTI feeds), evolving SENTINEL into a fully Autonomous Enterprise SOC Command Center.

Fifth, addressing state storage stability risks, production deployments will mandate SQLite configurations in `synchronous=NORMAL` mode (strictly avoiding `WAL` mode to eliminate cross-UID permission lock crashes in Docker environments), combined with latency-based Redis consumer group backpressure monitoring to guarantee 24/7 continuous operational availability for enterprise SOCs.

Concluding a rigorous research endeavor, while opening inspiring horizons in securing the digital realm through the power of Agentic Artificial Intelligence.

References

- [1] I. Sharafaldin, A. H. Lashkari, and A. A. Ghorbani, “Toward Generating a New Intrusion Detection Dataset and Intrusion Traffic Characterization,” in *Proceedings of the 4th International Conference on Information Systems Security and Privacy (ICISSP)*, 2018, pp. 108–116.
- [2] G. González-Granadillo, S. González-Zarzosa, and R. Diaz, “Security Information and Event Management (SIEM): Analysis, Trends, and Usage in Critical Infrastructures,” *Sensors*, vol. 21, no. 14, art. 4759, 2021.
- [3] R. Zuech, T. M. Khoshgoftaar, and R. Wald, “Intrusion Detection and Big Heterogeneous Data: A Survey,” *Journal of Big Data*, vol. 2, art. 3, 2015.
- [4] S. Tariq, M. Baruwat Chhetri, S. Nepal, and C. Paris, “Alert Fatigue in Security Operations Centres: Research Challenges and Opportunities,” *ACM Computing Surveys*, vol. 57, no. 9, art. 224, 2025.
- [5] Ismail, R. Kurnia, Z. A. Brata, G. A. Nelistiani, S. Heo, H. Kim, and H. Kim, “Toward Robust Security Orchestration and Automated Response in Security Operations Centers with a Hyper-Automation Approach Using Agentic Artificial Intelligence,” *Information*, vol. 16, no. 5, art. 365, 2025.
- [6] B. P. Welford, “Note on a Method for Calculating Corrected Sums of Squares and Products,” *Technometrics*, vol. 4, no. 3, pp. 419–420, 1962.
- [7] P. Lewis *et al.*, “Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks,” *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 33, pp. 9459–9474, 2020.
- [8] LangChain AI, “LangGraph: Multi-Actor, Stateful LLM Applications,” *GitHub repository*, 2024. [Online]. Available: <https://github.com/langchain-ai/langgraph>
- [9] K. Greshake, S. Abdelnabi, S. Mishra, A. Endres, T. Holz, and M. Fritz, “More than you’ve asked for: A Comprehensive Analysis of Novel Prompt Injection Threats to Application-Integrated Large Language Models,” *arXiv preprint arXiv:2302.12173*, 2023.
- [10] P. Cichonski, T. Millar, T. Grance, and K. Scarfone, “Computer Security Incident Handling Guide,” NIST Special Publication 800-61 Revision 2, 2012.
- [11] G. V. Cormack, C. L. A. Clarke, and S. Buettcher, “Reciprocal Rank Fusion Outperforms Condorcet and Individual Bootstrap Product Algorithms,” in *Proceedings of the 32nd International ACM SIGIR Conference on Research and Development in Information Retrieval*, 2009.

- [12] R. Pandey and A. Bhujang, “Poisoning the Watchtower: Prompt Injection Attacks Against LLM-Augmented Security Operations Through Adversarial Log Content,” *arXiv preprint arXiv:2605.24421*, 2026.
- [13] F. Blefari, C. Cosentino, F. A. Pironti, A. Furfaro, and F. Marozzo, “CyberRAG: An Agentic RAG cyber attack classification and reporting tool,” *Future Generation Computer Systems*, vol. 176, art. 108186, 2026.
- [14] Z. Liu and A. Anwar, “AutoBnB-RAG: Enhancing Multi-Agent Incident Response with Retrieval-Augmented Generation,” *arXiv preprint arXiv:2508.13118*, 2025.
- [15] J. Zhang *et al.*, “When LLMs Meet Cybersecurity: A Systematic Literature Review,” *Cybersecurity*, vol. 8, 2025.
- [16] A. Abdennebi, N. Kara, L. Lahlou, and H. Ould-Slimane, “LanG – A Governance-Aware Agentic AI Platform for Unified Security Operations,” *arXiv preprint arXiv:2604.05440*, 2026.
- [17] R. Sahay *et al.*, “Policy-Guided Threat Hunting: An LLM Enabled Framework with Splunk SOC Triage,” *arXiv preprint arXiv:2603.23966*, 2026.
- [18] A. Vaswani *et al.*, “Attention is All You Need,” in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 30, pp. 5998–6008, 2017.
- [19] A. Gholami *et al.*, “A Survey of Quantization Methods for Efficient Neural Network Inference,” *arXiv preprint arXiv:2103.13630*, 2021.
- [20] J. Johnson, M. Douze, and H. Jégou, “Billion-Scale Similarity Search with GPUs,” *IEEE Transactions on Big Data*, vol. 7, no. 3, pp. 535–547, 2019.
- [21] S. Robertson and H. Zaragoza, “The Probabilistic Relevance Framework: BM25 and Beyond,” *Foundations and Trends in Information Retrieval*, vol. 3, no. 4, pp. 333–380, 2009.
- [22] P. He, J. Zhu, S. He, J. Li, and M. R. Lyu, “Drain: An Online Log Parsing Approach with Fixed Depth Tree,” in *Proceedings of the IEEE International Conference on Web Services (ICWS)*, 2017, pp. 33–40.
- [23] Gemma Team, Google, “Gemma 2: Improving Open Language Models at a Practical Size,” *arXiv preprint arXiv:2408.00118*, 2024.
- [24] B. E. Strom, A. Applebaum, D. P. Miller, K. C. Nickels, A. G. Pennington, and C. B. Thomas, “MITRE ATT&CK: Design and Philosophy,” The MITRE Corporation, Technical Report MP180360, 2018.
- [25] N. Reimers and I. Gurevych, “Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks,” in *Proceedings of the Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2019.
- [26] L. Zheng *et al.*, “Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena,” in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 36, 2023.

- [27] S. Myneni, A. Chowdhary, A. Sabur, S. Sengupta, G. Agrawal, D. Huang, and M. Kang, “DAPT 2020 – Constructing a Benchmark Dataset for Advanced Persistent Threats,” in *Deployable Machine Learning for Security Defense (MLHat 2020)*, Communications in Computer and Information Science, vol. 1271, Springer, 2020, pp. 138–163.
- [28] S. Es, J. James, L. Espinosa-Anke, and S. Schockaert, “RAGAS: Automated Evaluation of Retrieval Augmented Generation,” *arXiv preprint arXiv:2309.15217*, 2023.
- [29] H. Krawczyk, M. Bellare, and R. Canetti, “HMAC: Keyed-Hashing for Message Authentication,” Internet Engineering Task Force, RFC 2104, 1997.
- [30] OWASP GenAI Security Project, “OWASP Top 10 for Large Language Model Applications,” Version 2025, The OWASP Foundation, 2025. [Online]. Available: <https://genai.owasp.org/>
- [31] S. Rose, O. Borchert, S. Mitchell, and S. Connelly, “Zero Trust Architecture,” NIST Special Publication 800-207, National Institute of Standards and Technology, 2020.